

# THE PITFALLS OF ACHIEVING TAGGED CAUSAL DELIVERY

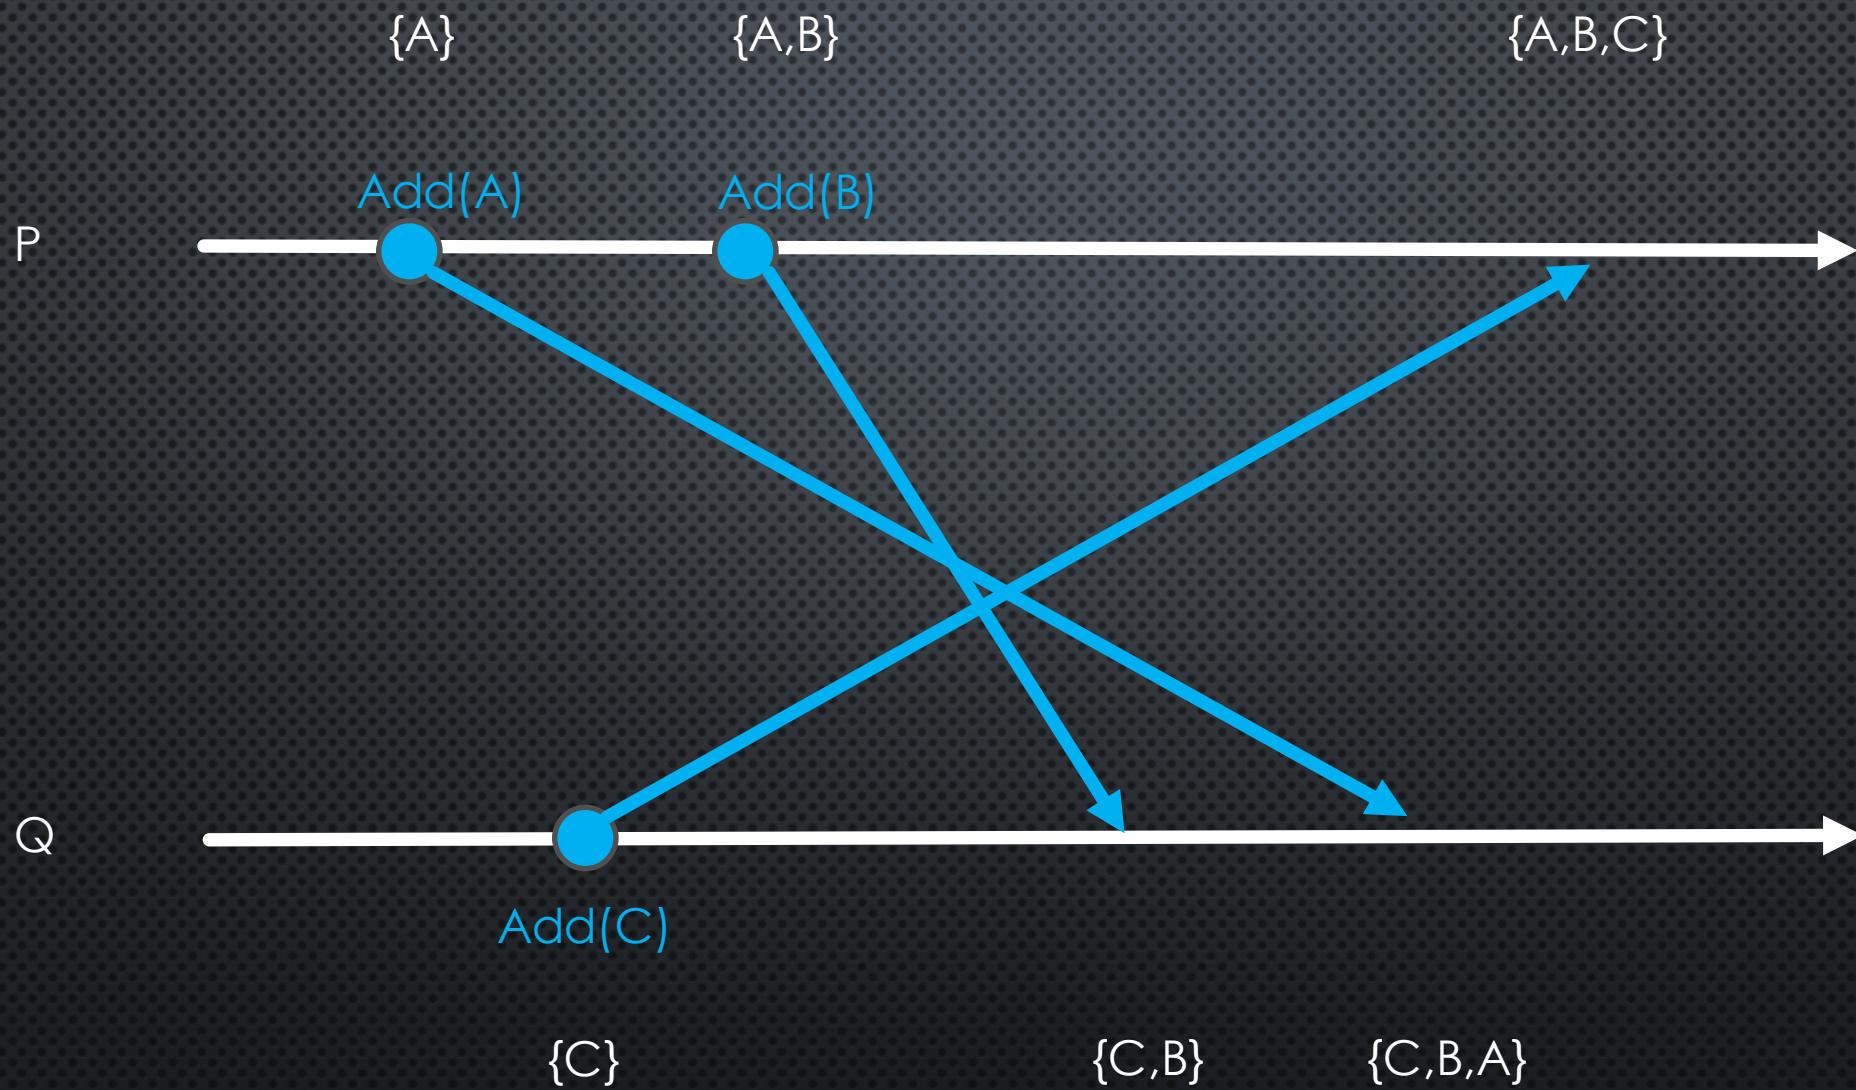
GEORGES YOUNES

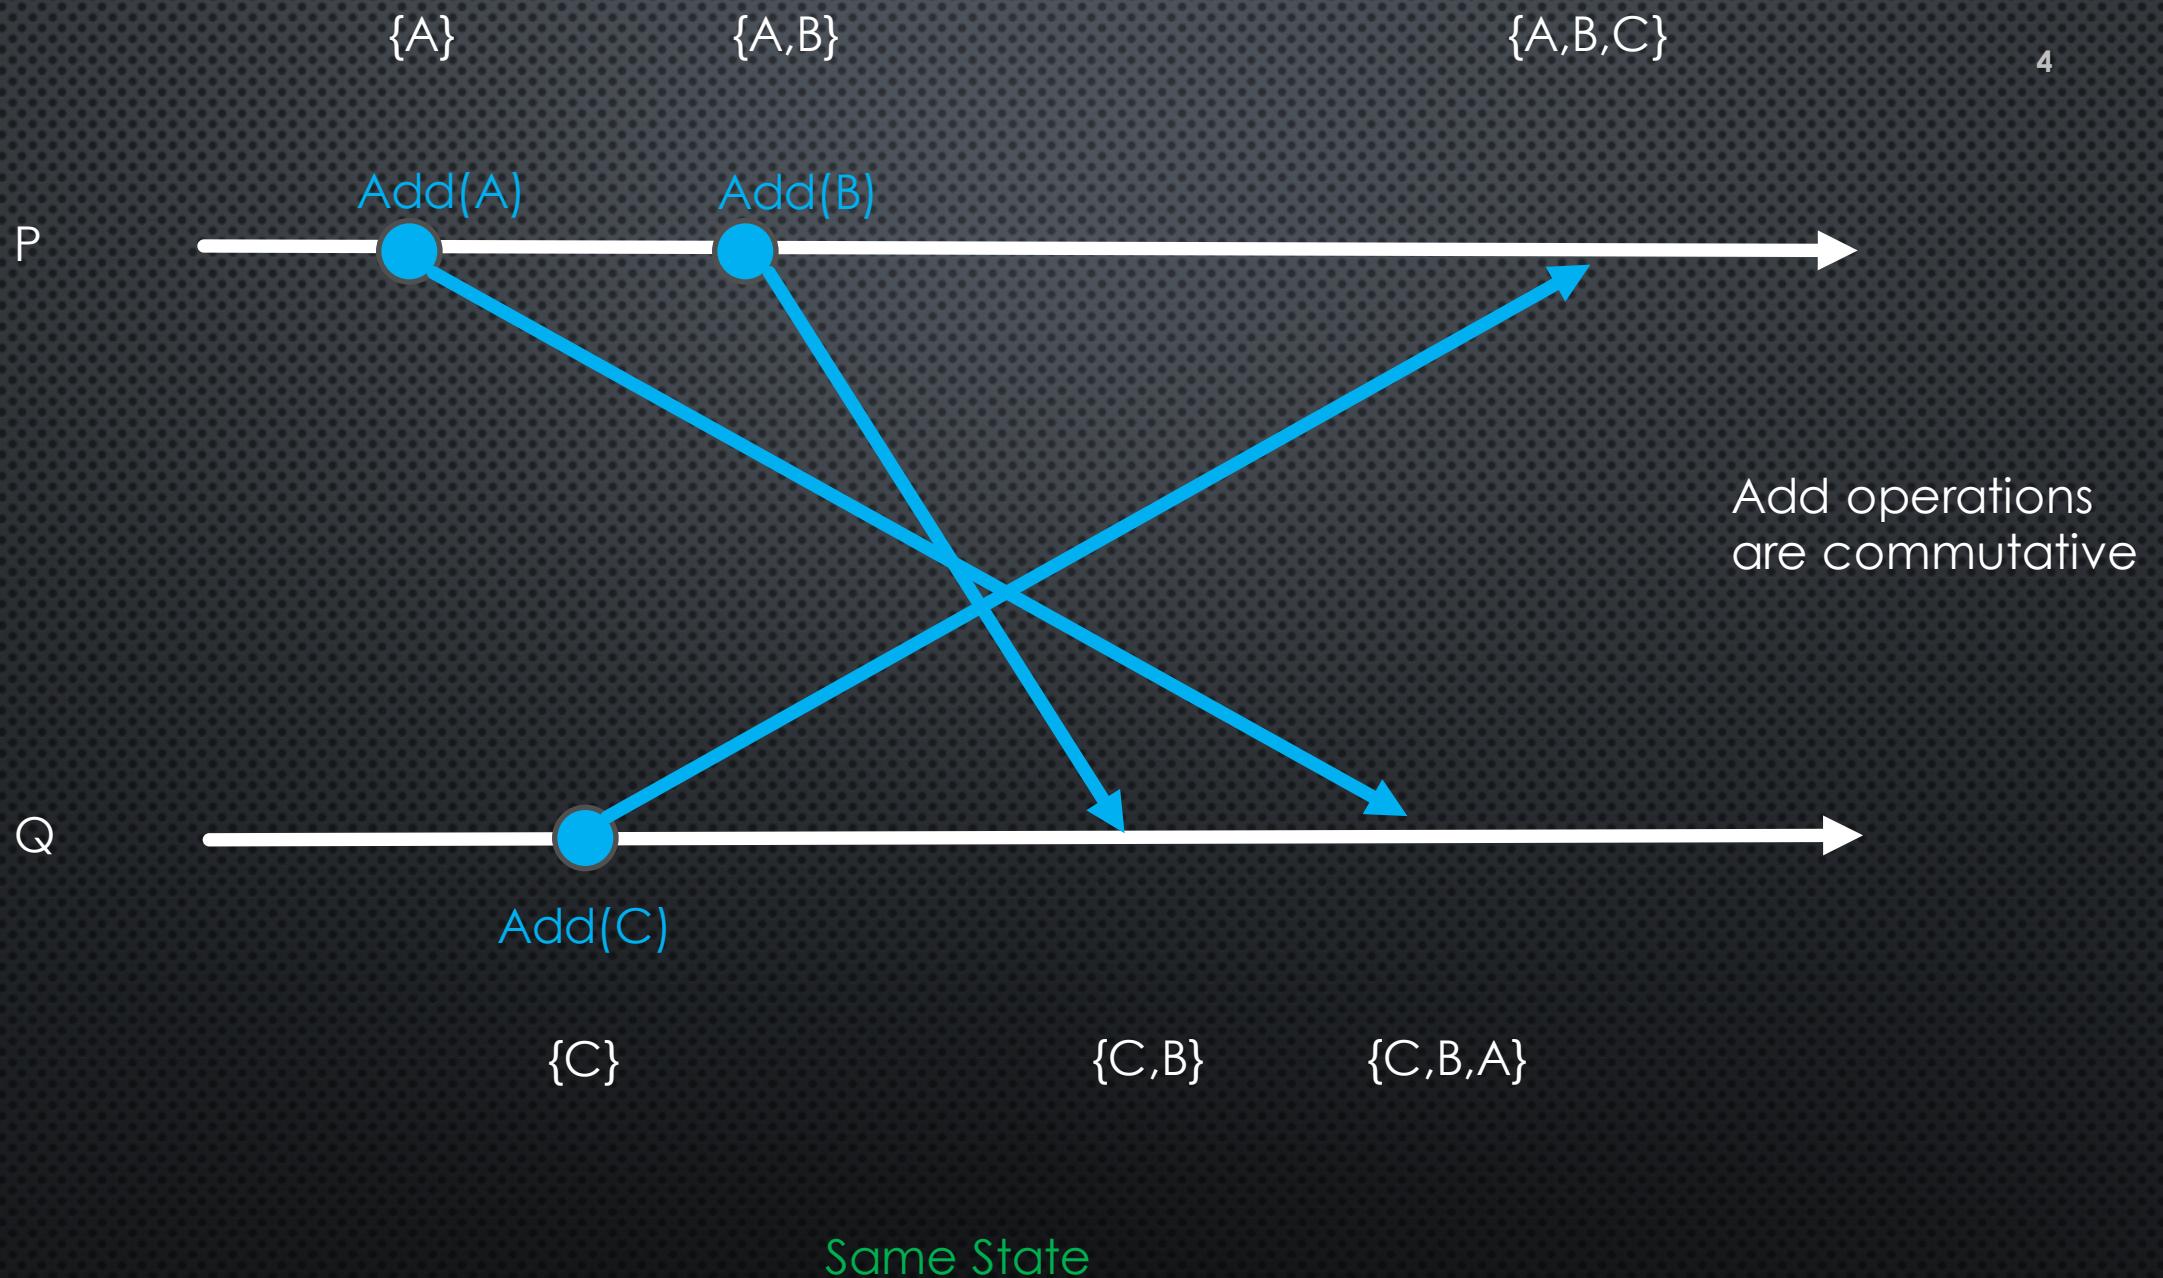
PAPOC'18

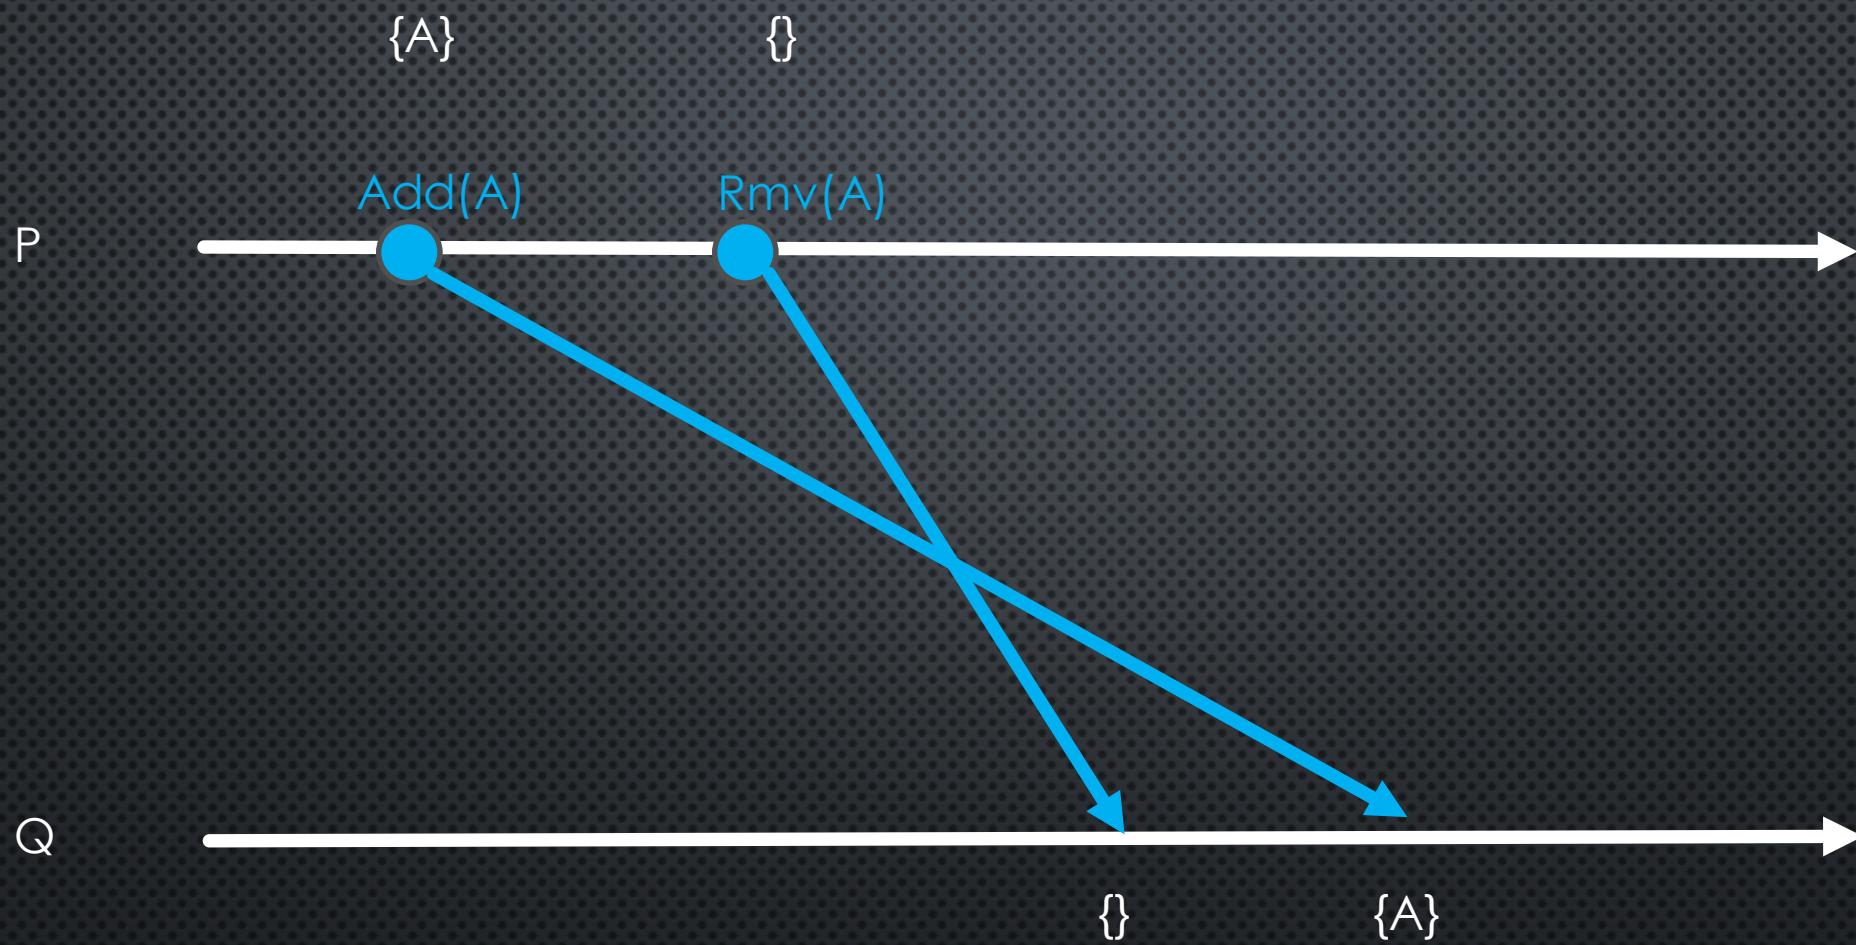
APRIL 23, 2018

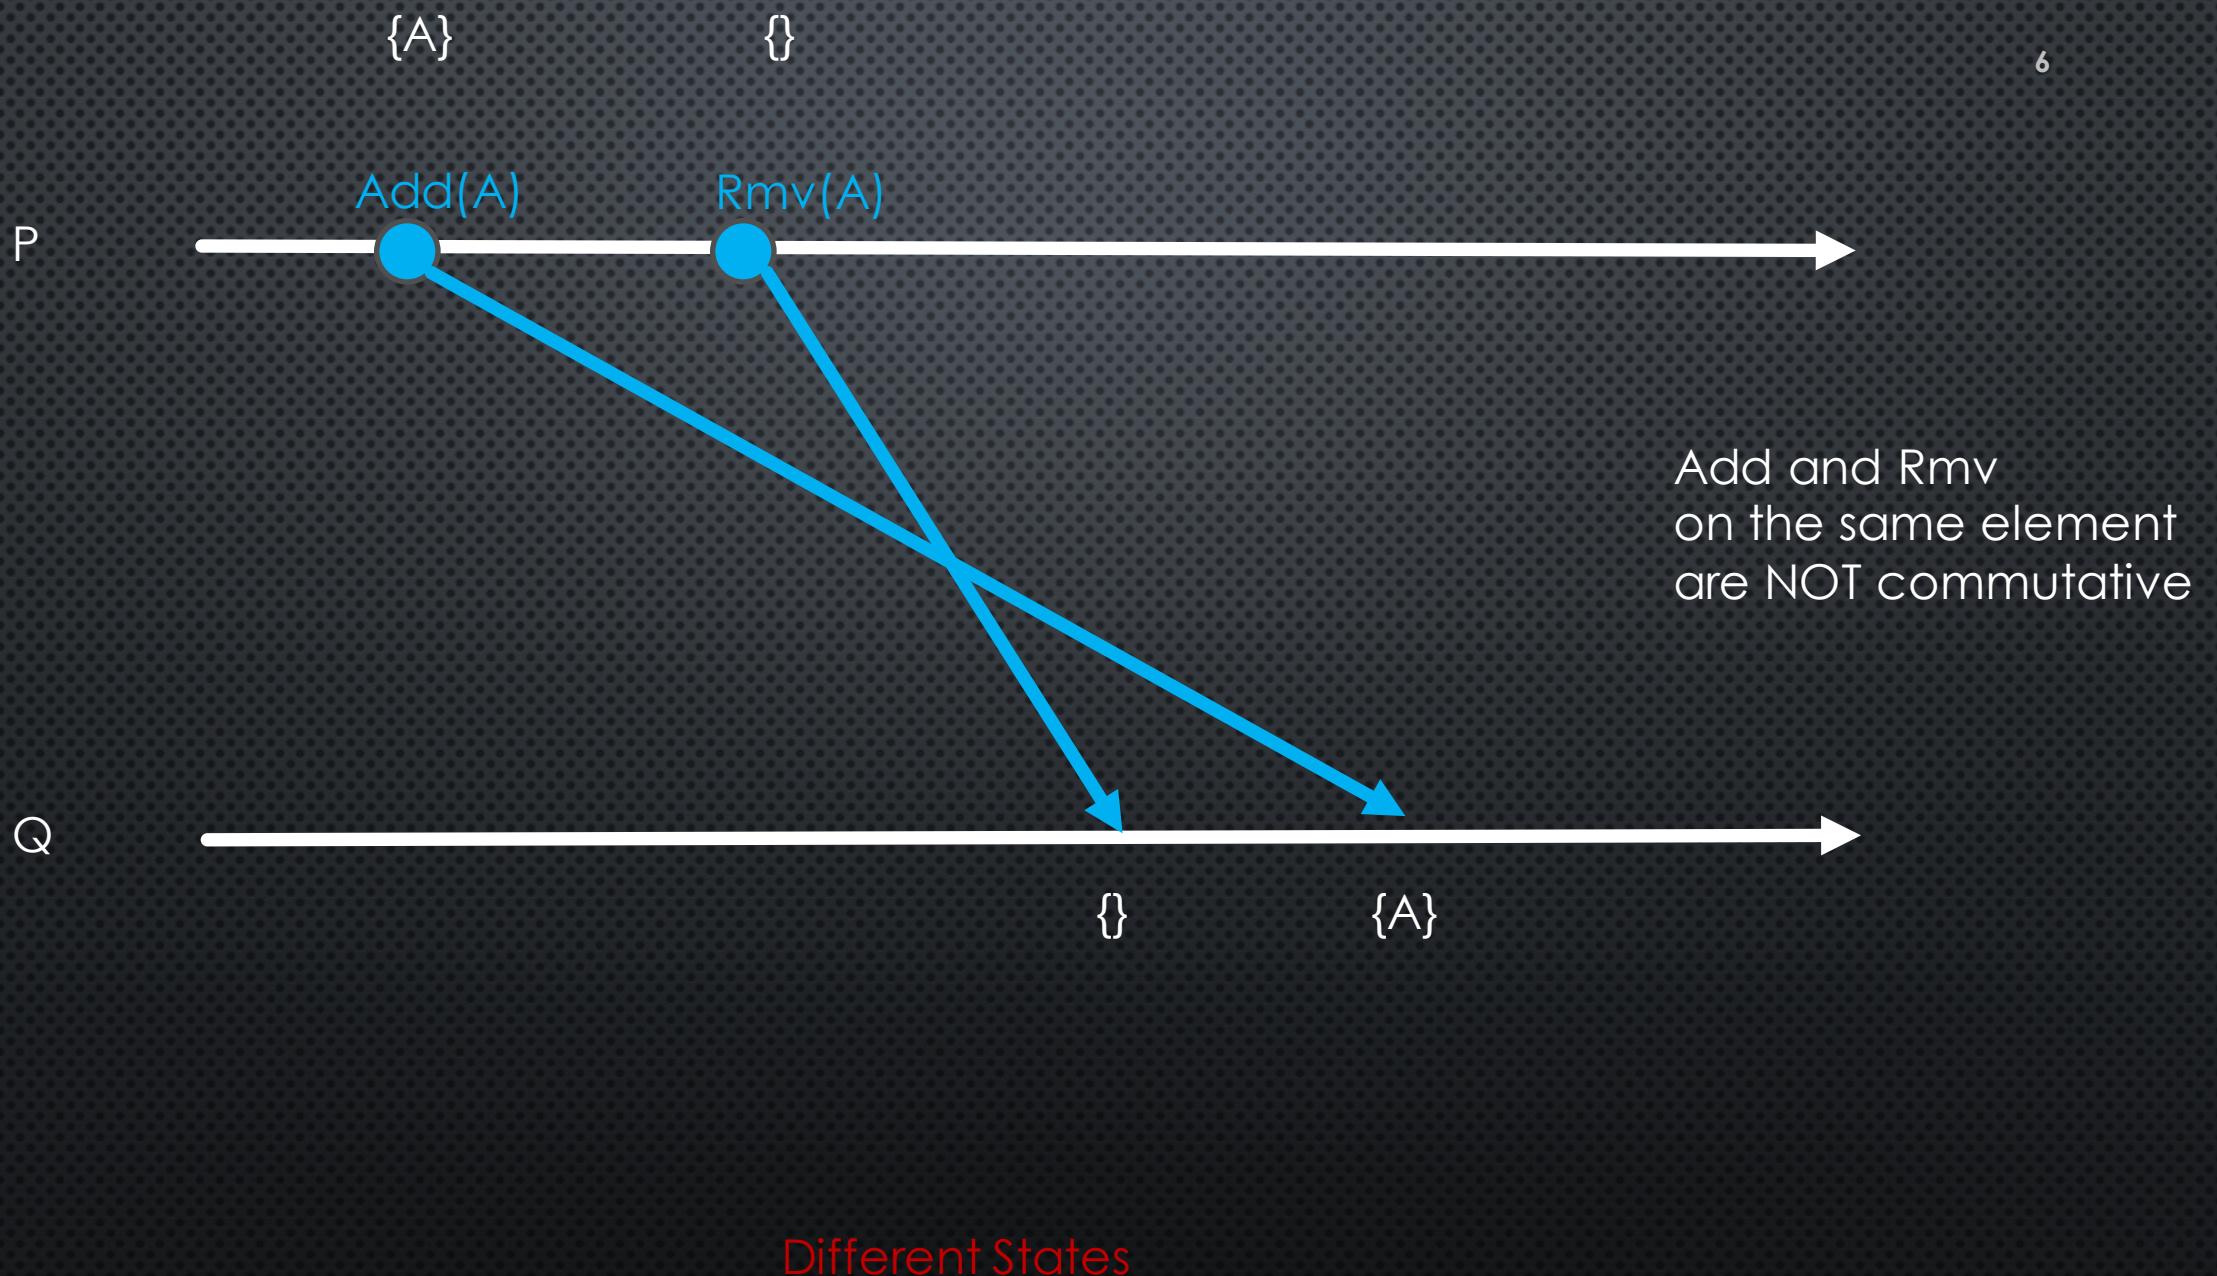
# USE CASE: REPLICATED SET

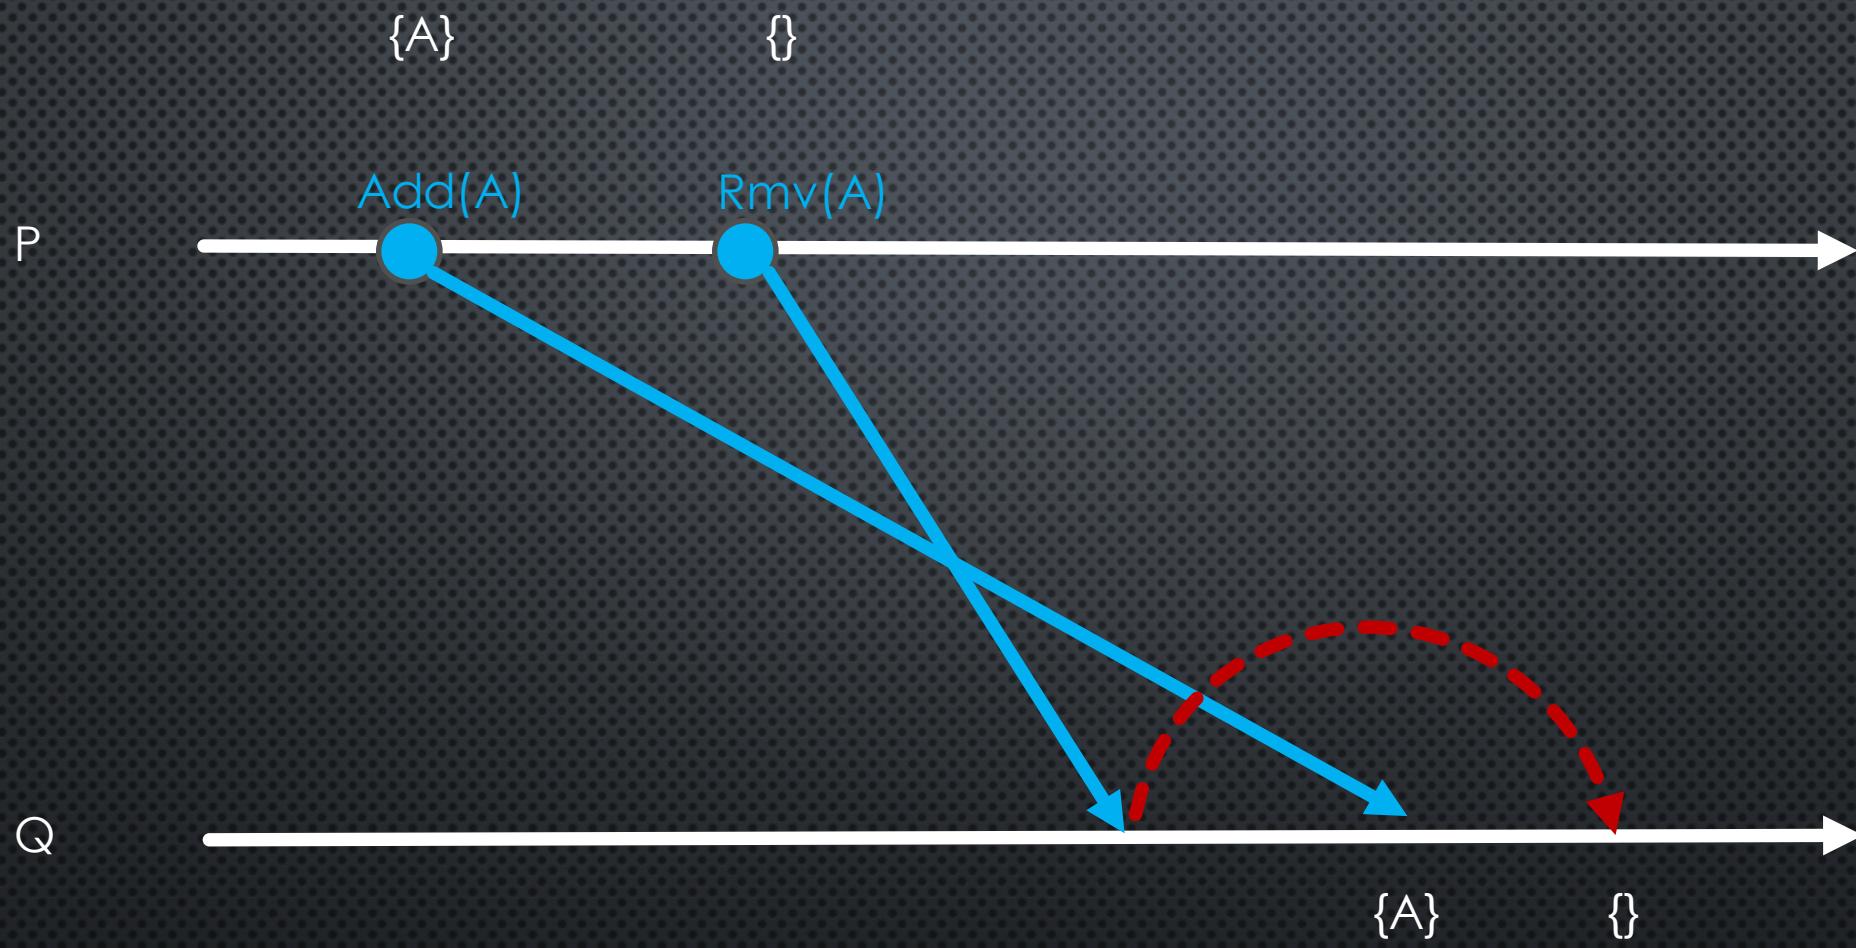
- API:
  - ADD(ELEMENT): ADDS THE ELEMENT TO THE SET
  - RMV(ELEMENT): REMOVES THE ELEMENT FROM THE SET
  - QUERY(): RETURNS ALL THE ELEMENTS IN THE SET





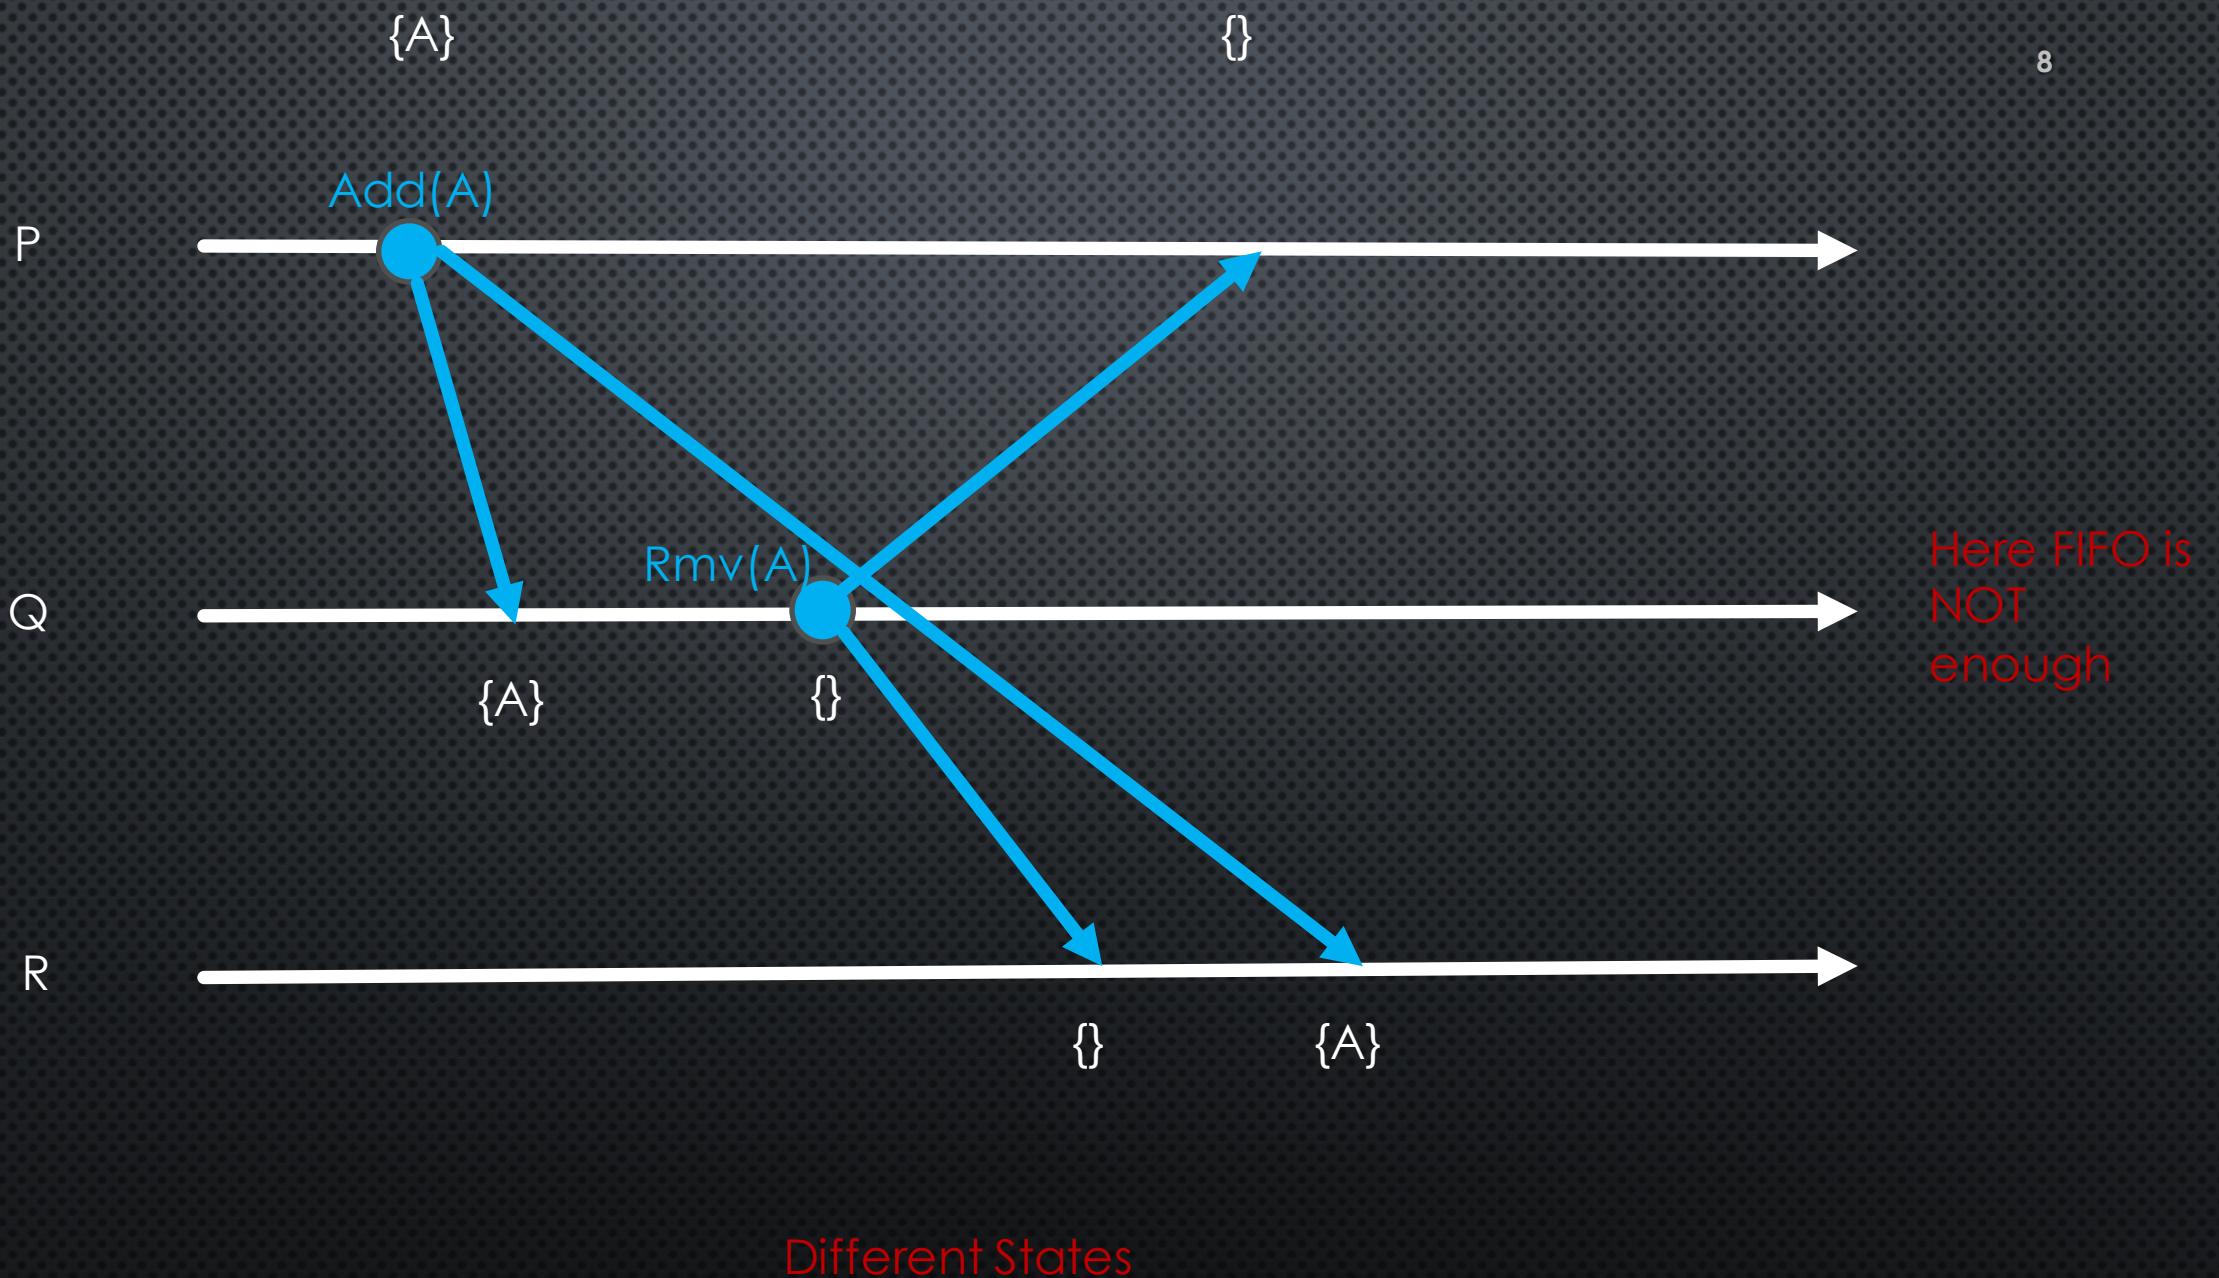


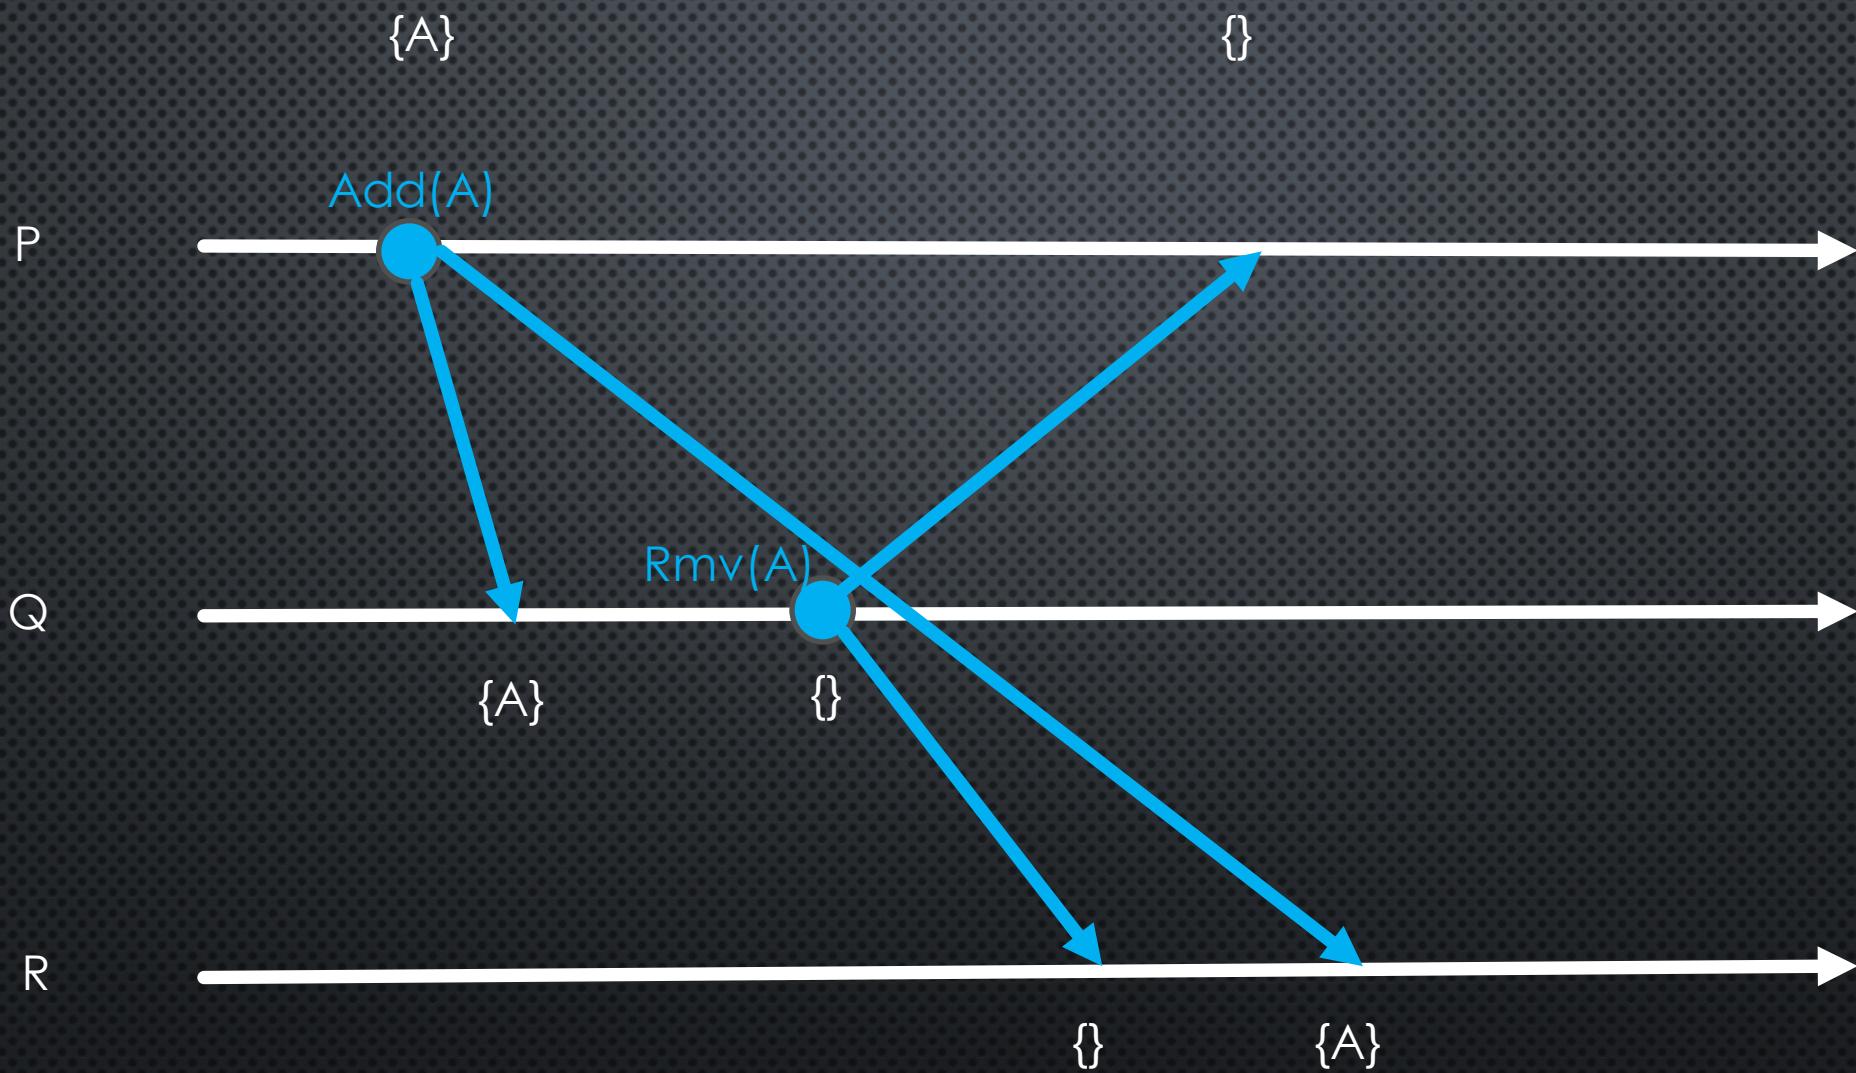


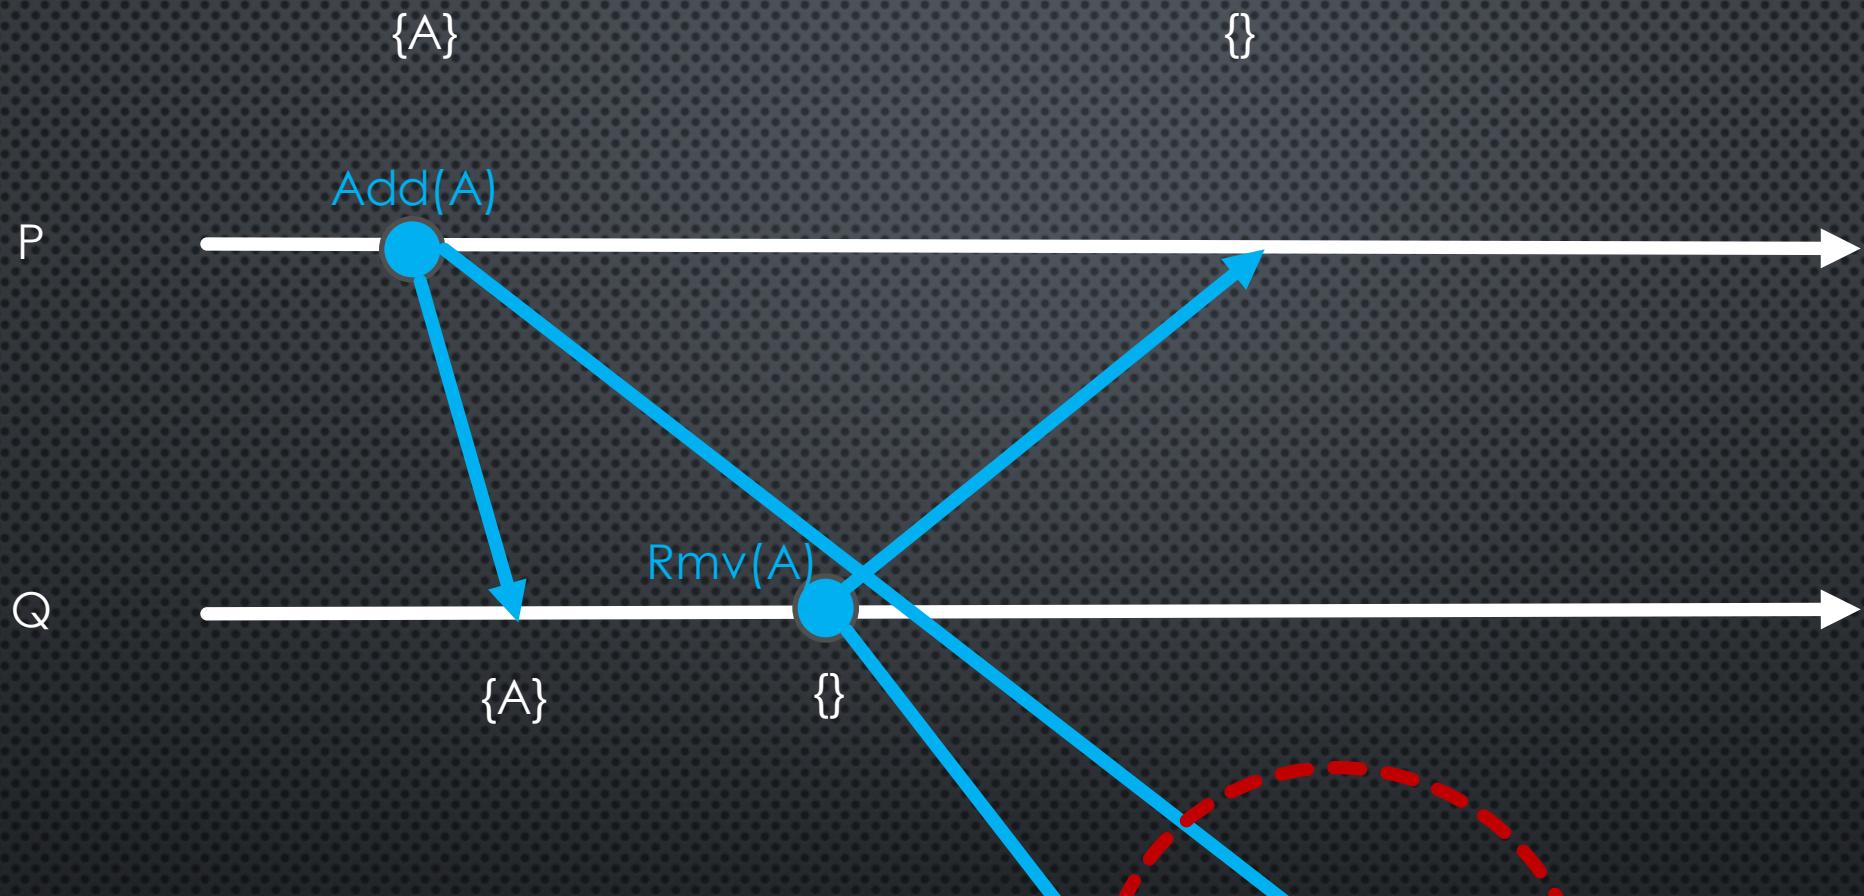


FIFO ordering abstraction

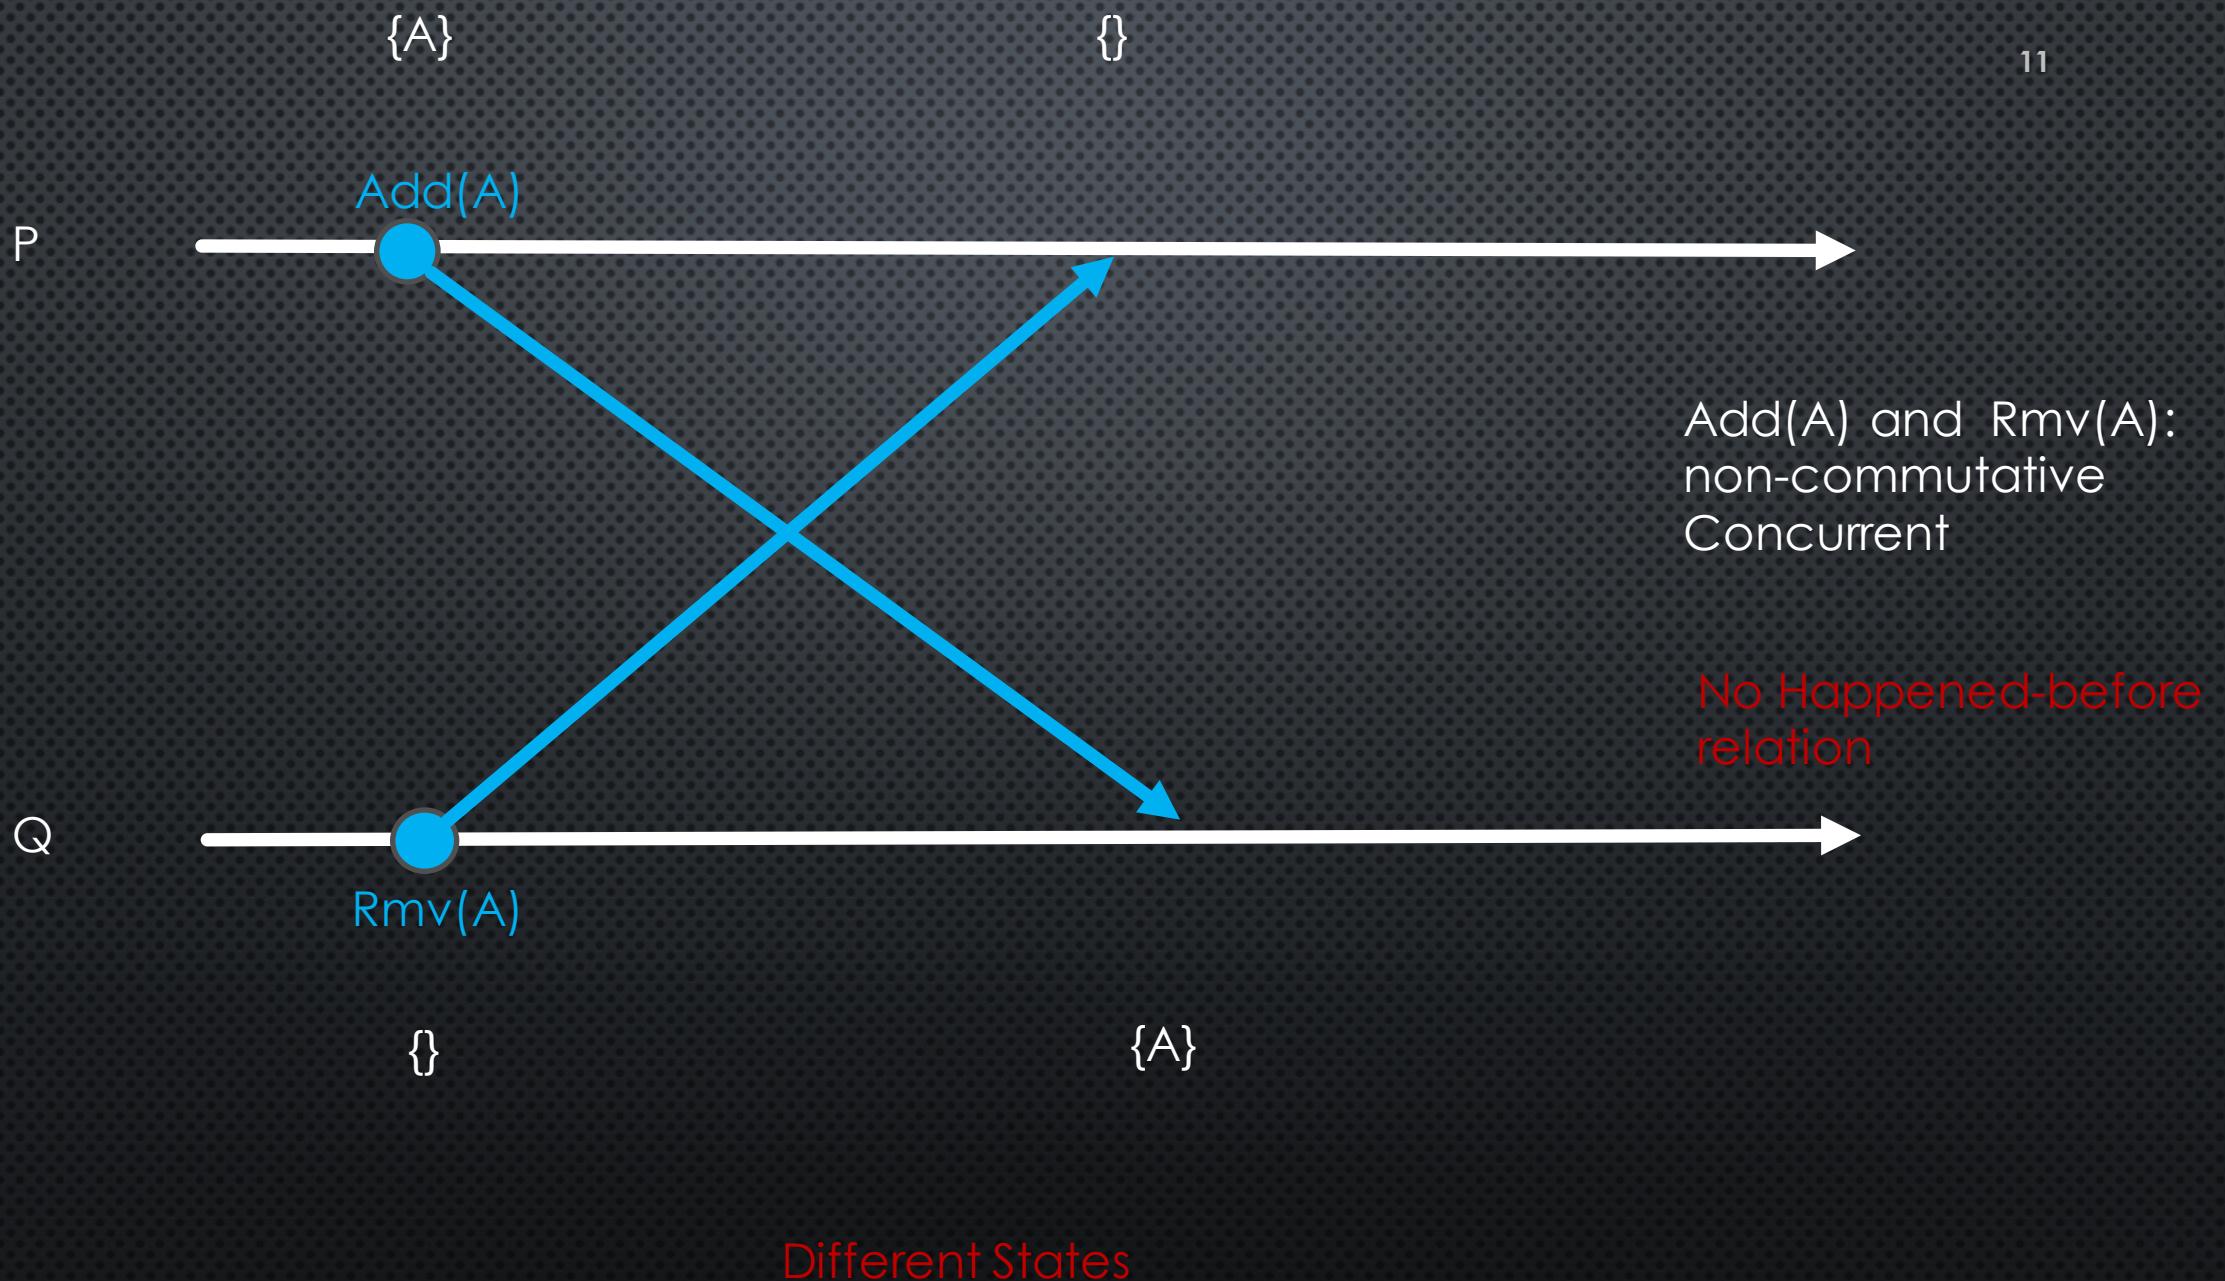
Same State

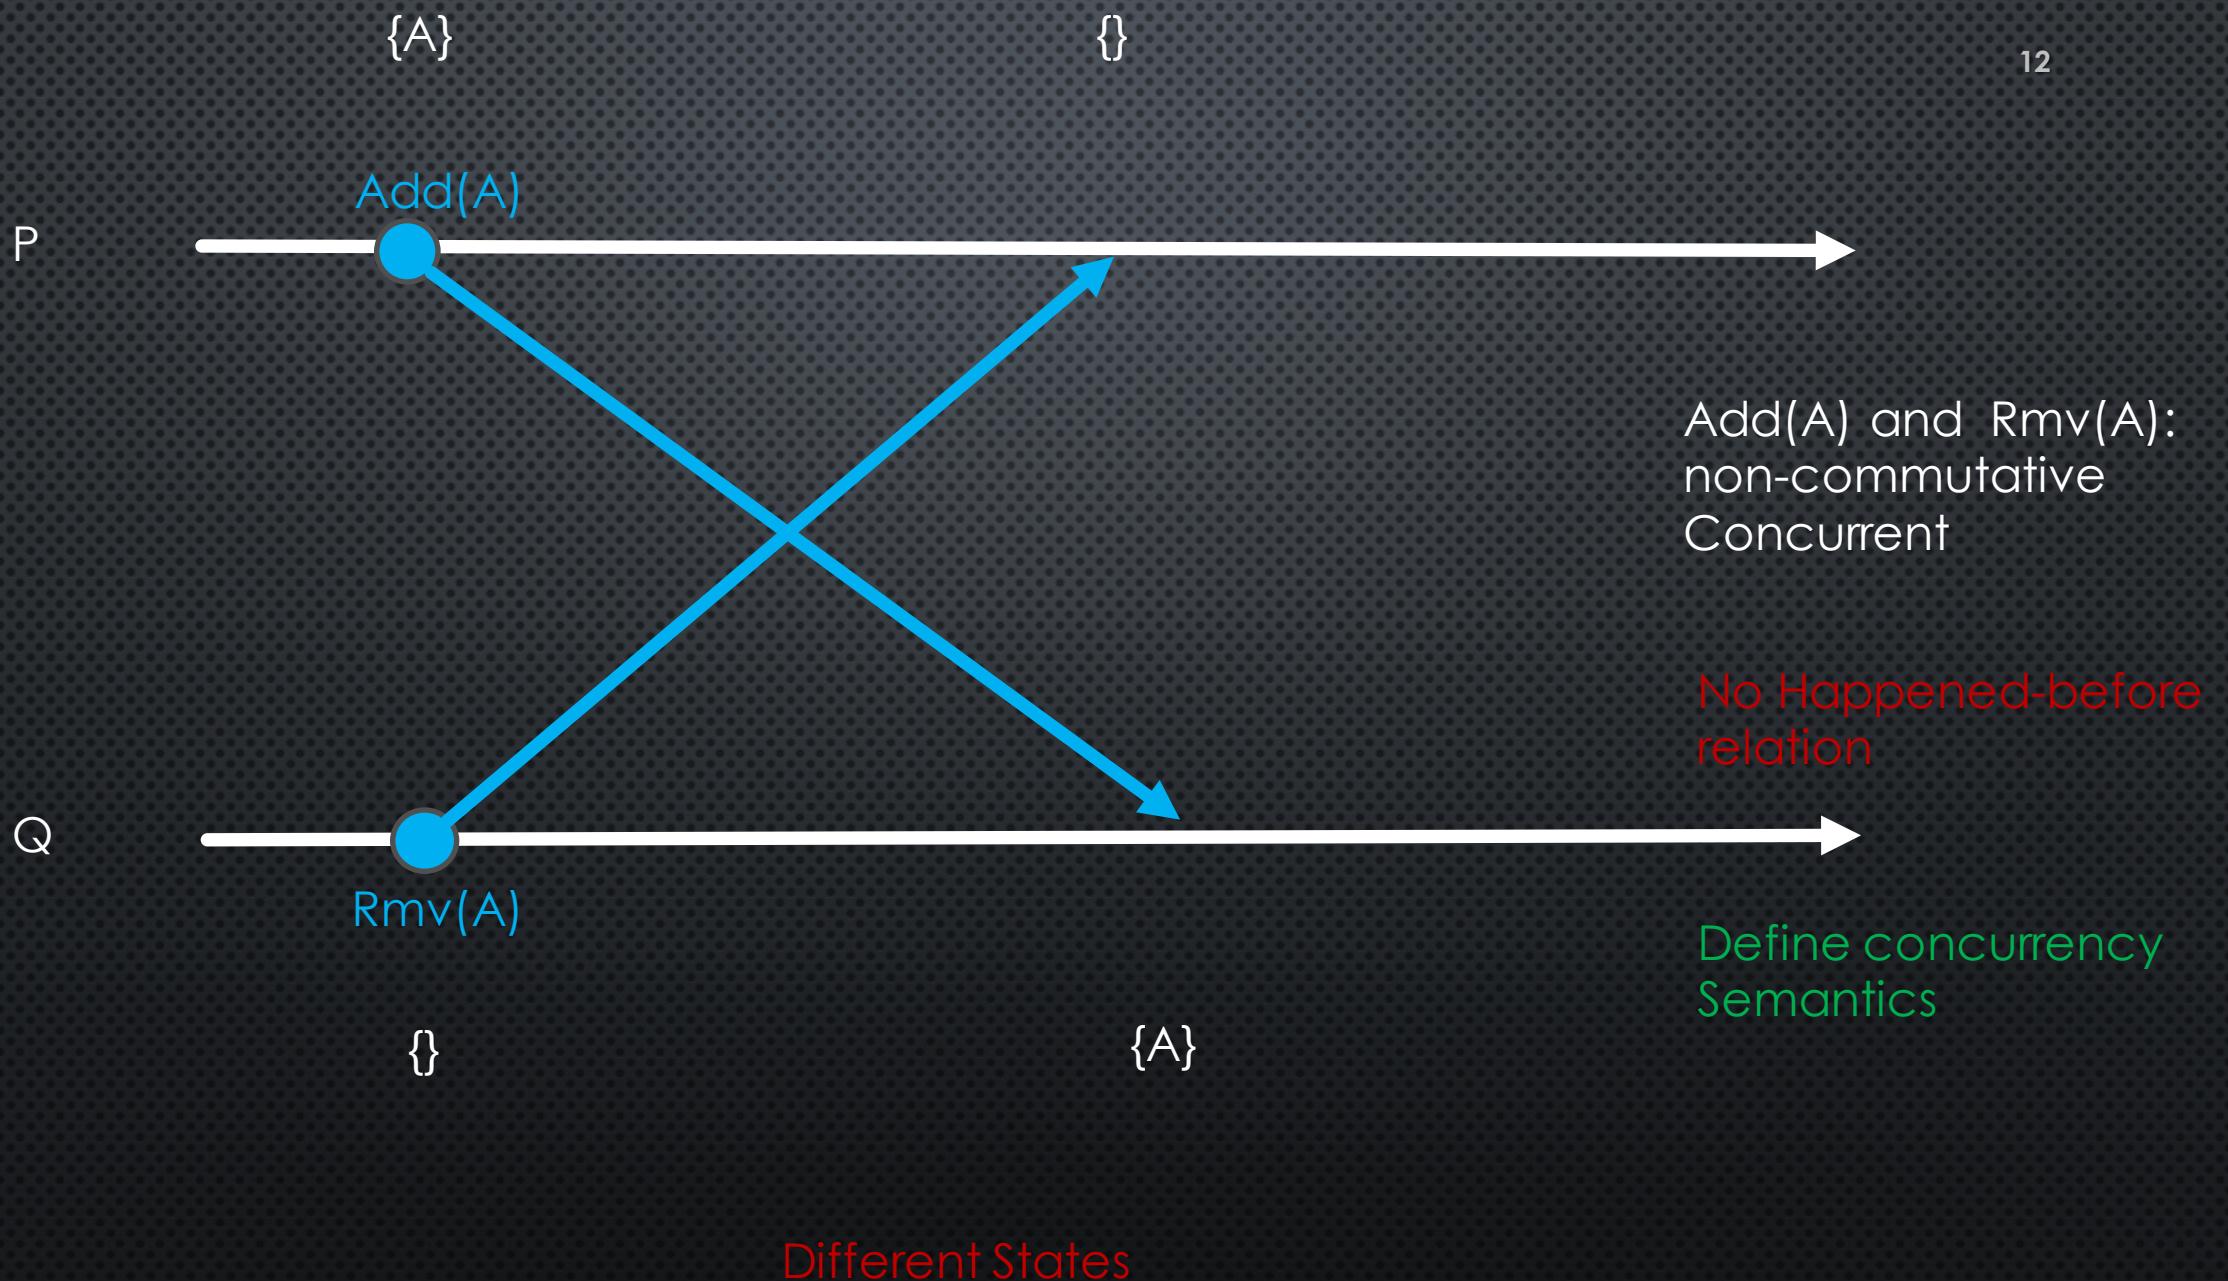






Causal ordering abstraction  
Same State





# SOLUTION 1

- WHAT YOU NEED:

- CAUSAL ORDER



Off-the-shelf middleware

- CONCURRENCY SEMANTICS



The “happened-before”  
relation between ops

# SOLUTION 1

- WHAT YOU NEED:

- CAUSAL ORDER



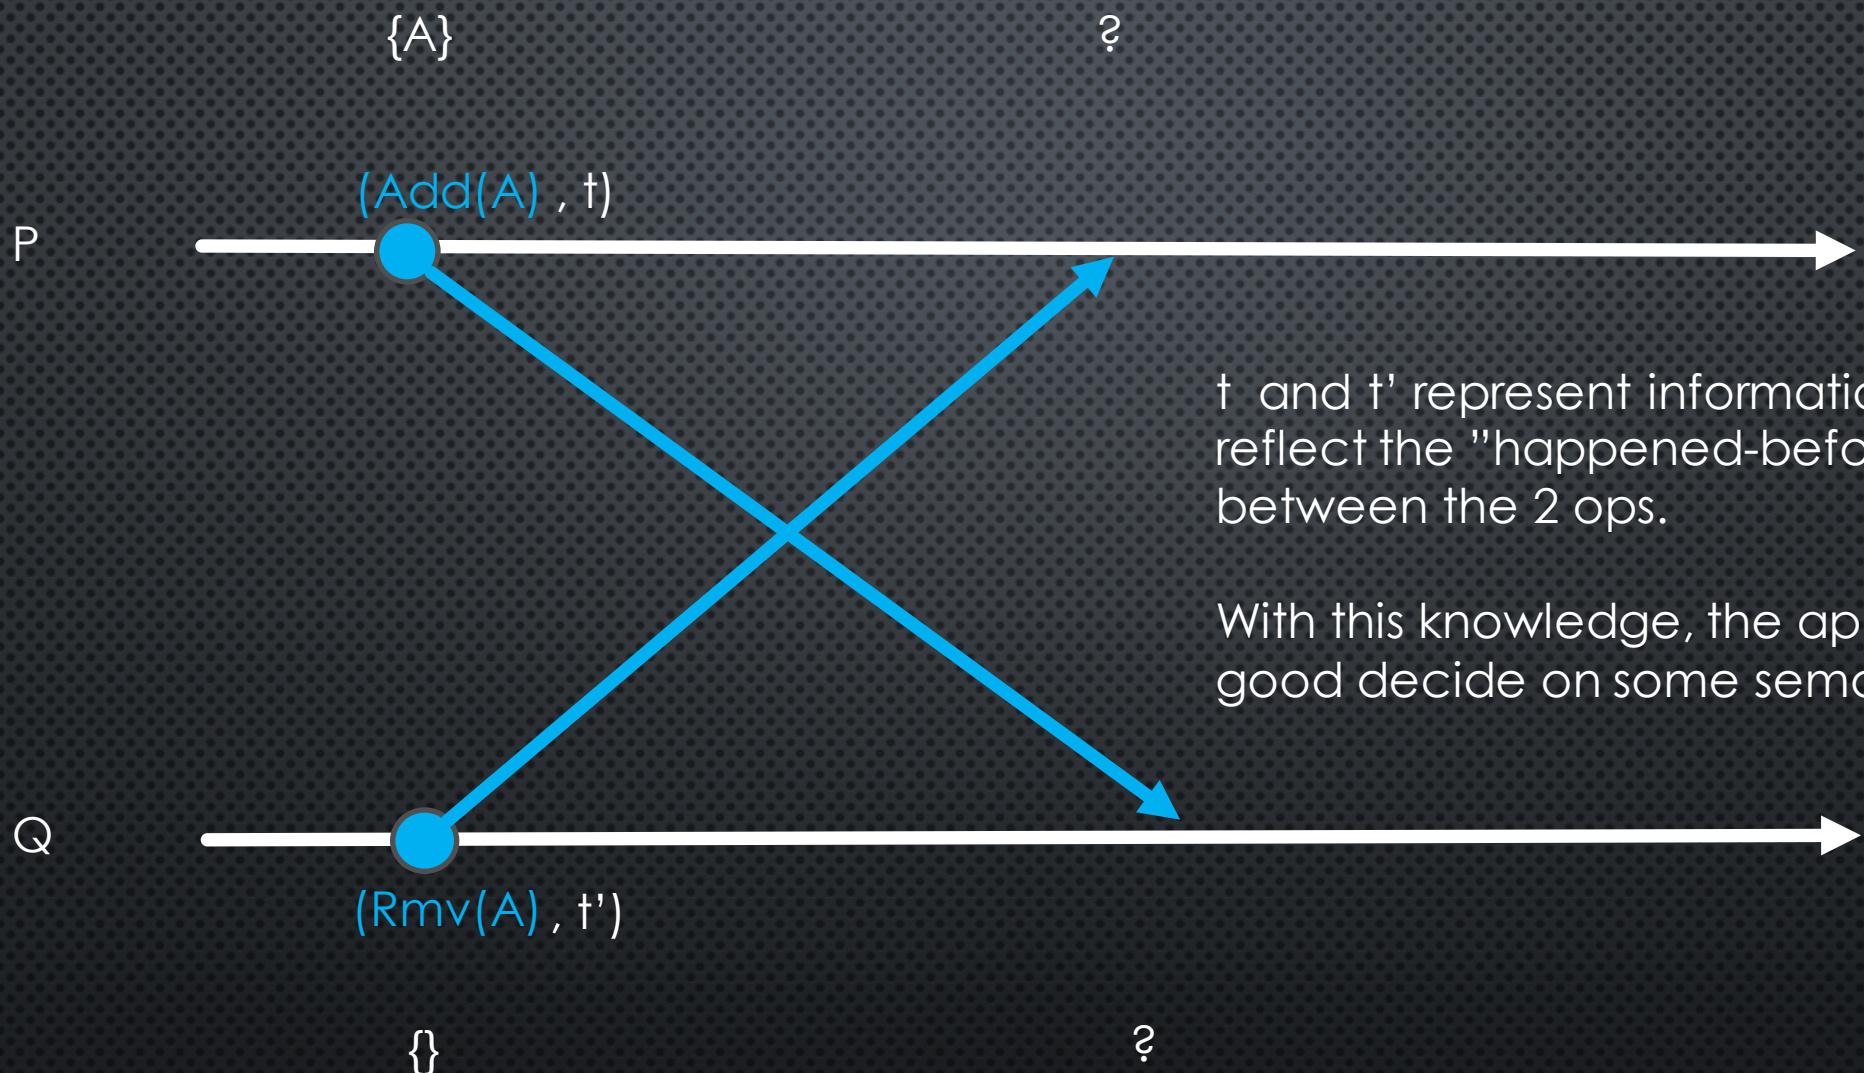
Off-the-shelf middleware

- CONCURRENCY SEMANTICS



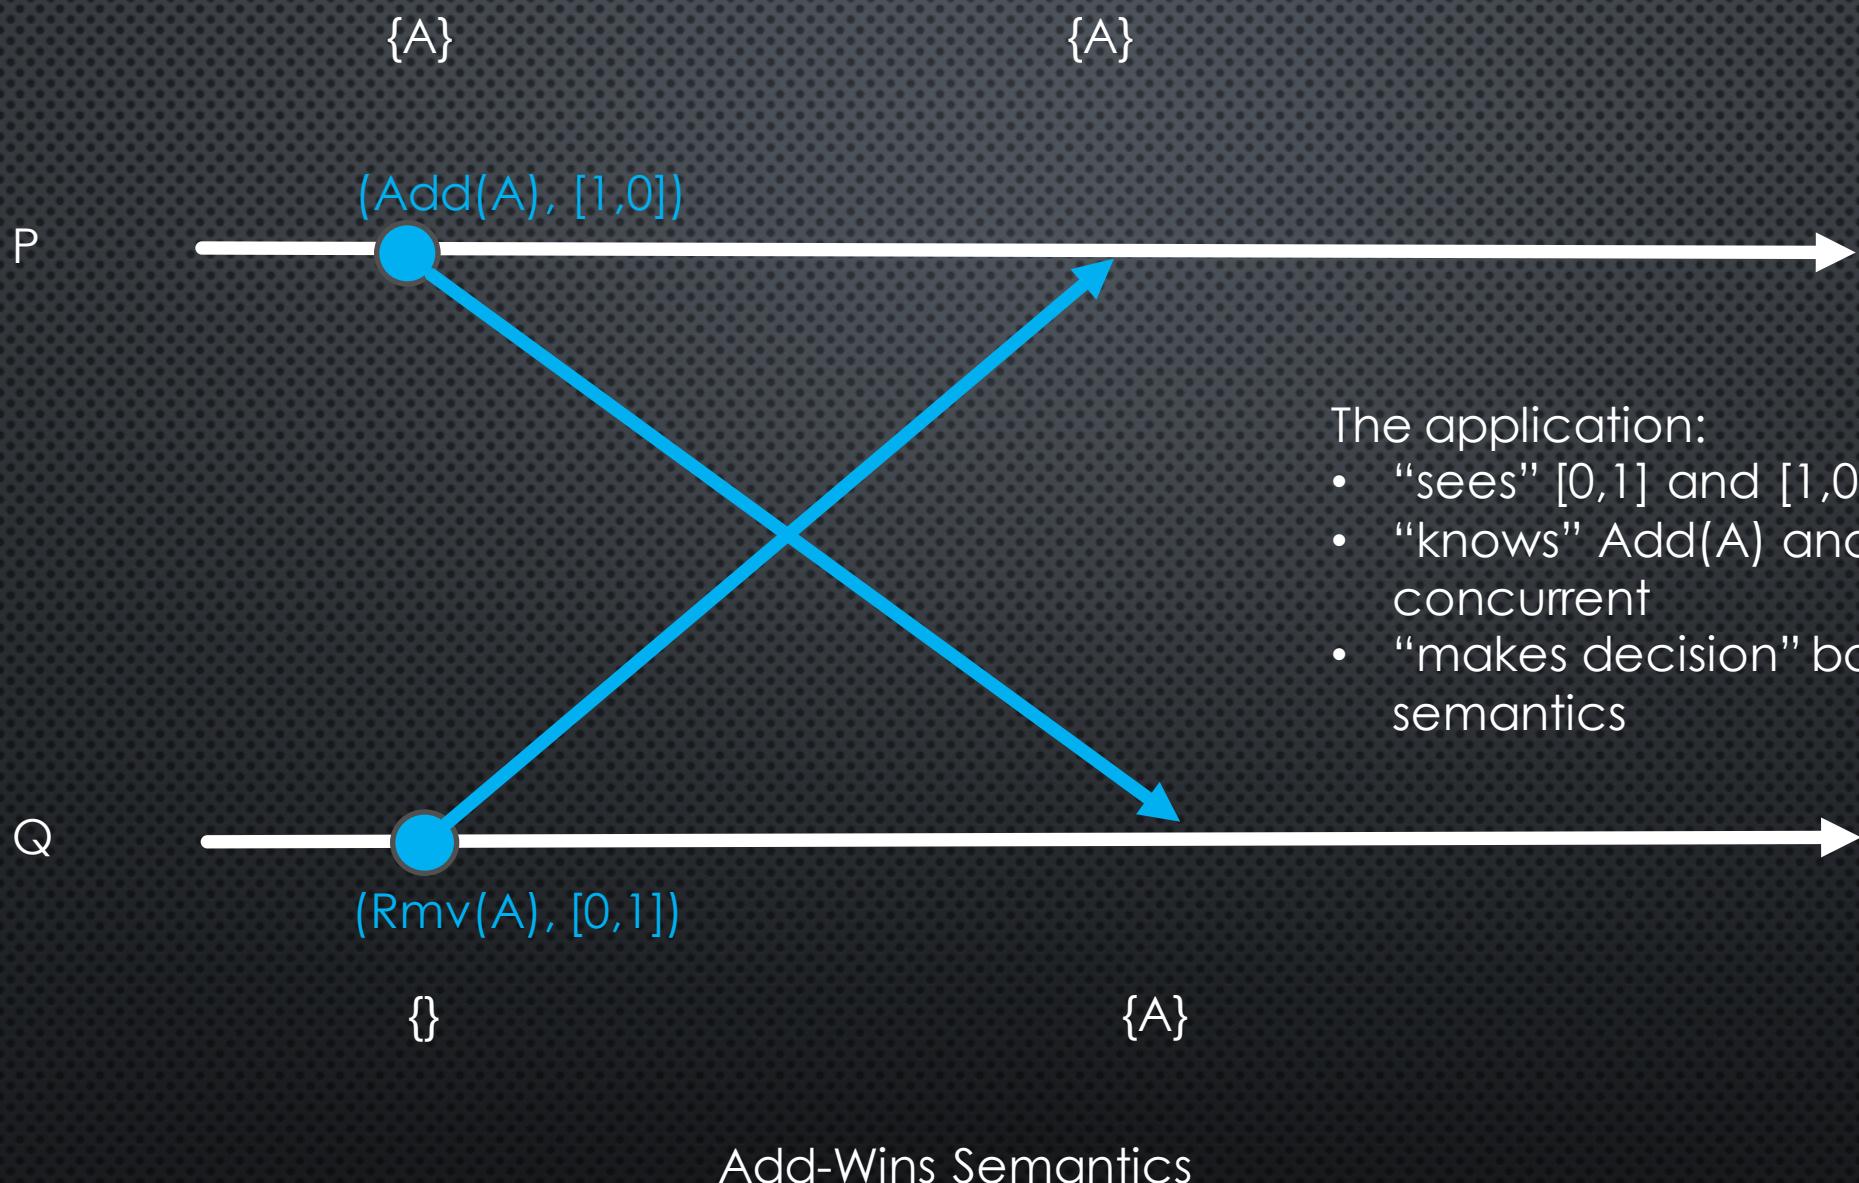
The “happened-before”  
relation between ops

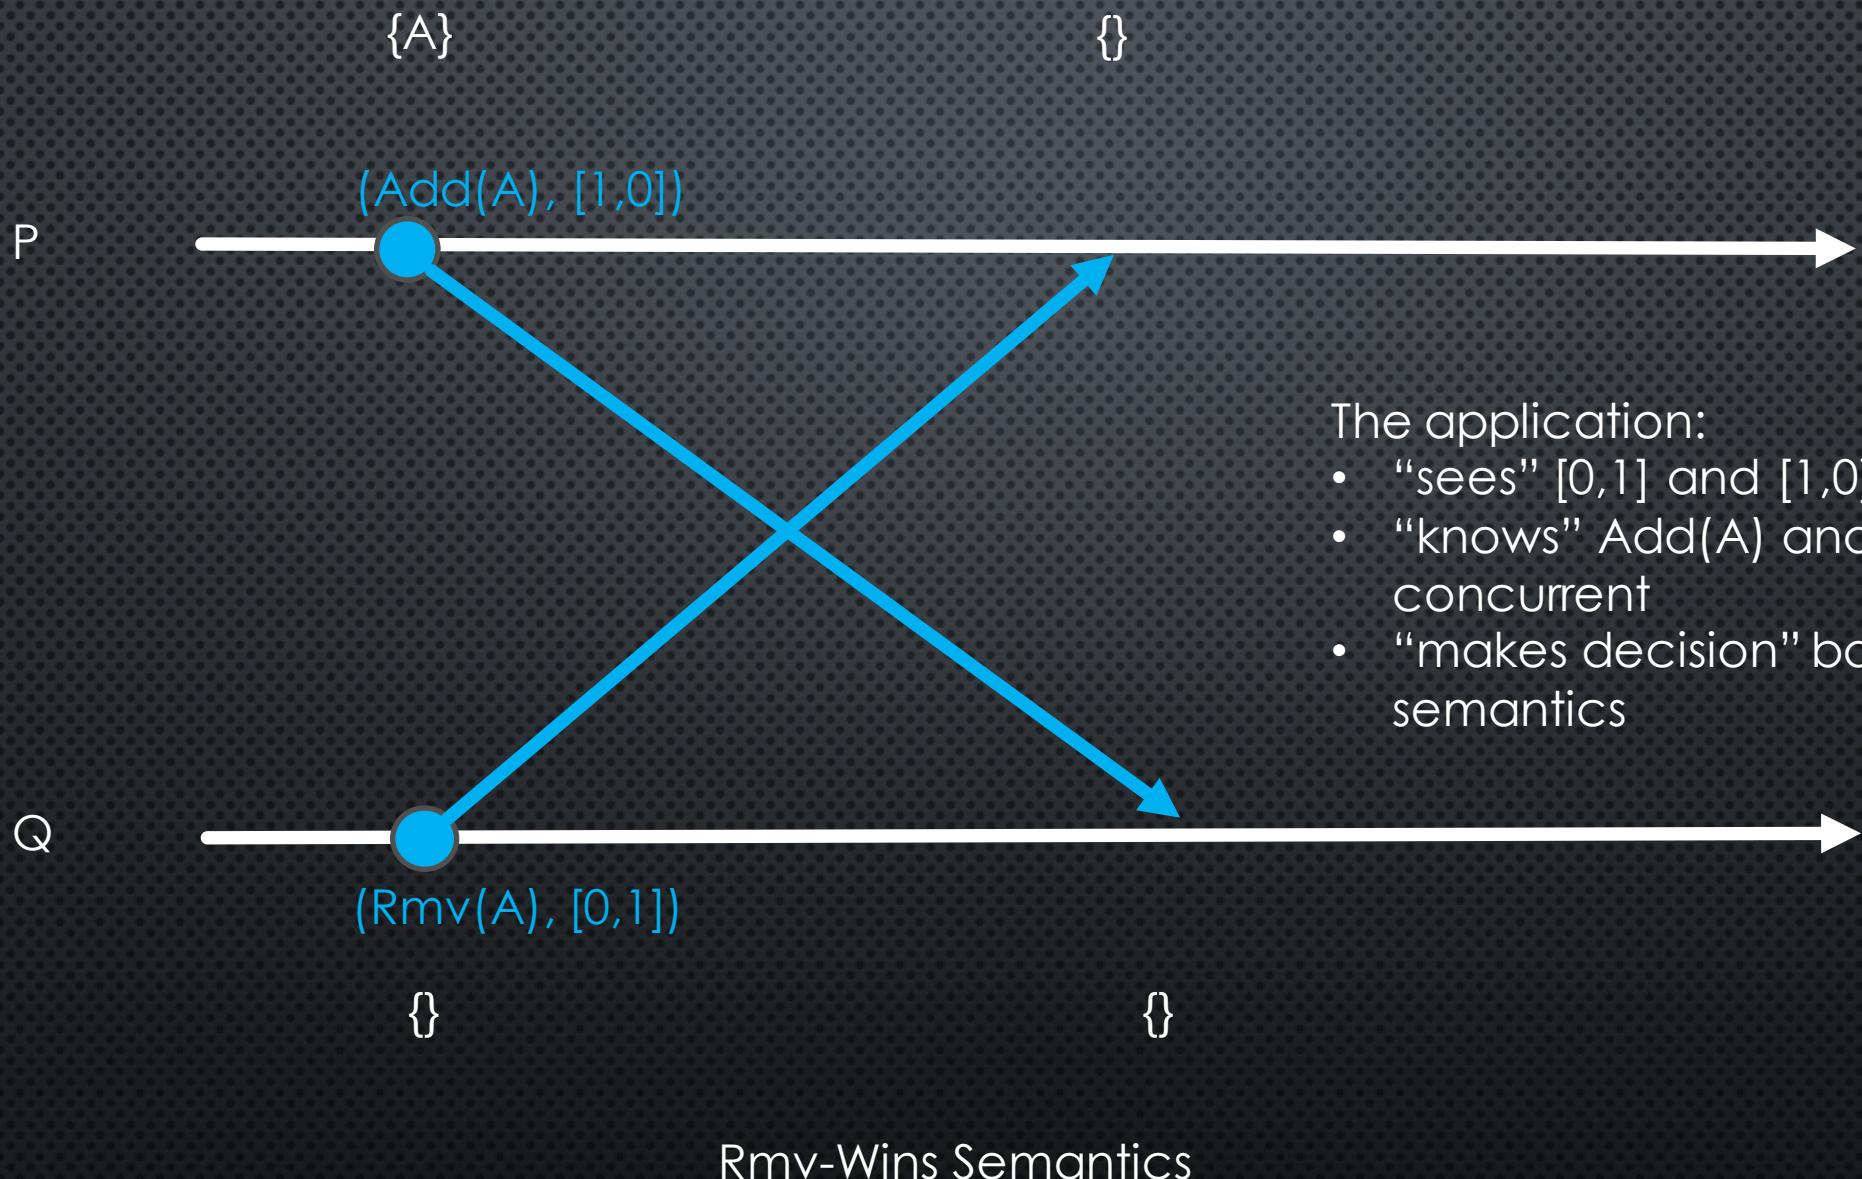
Requires explicit  
tagging of operations



$t$  and  $t'$  represent information that reflect the "happened-before" relation between the 2 ops.

With this knowledge, the application good decide on some semantics.





# SOLUTION 1

- WHAT YOU NEED:

- CAUSAL ORDER



Off-the-shelf middleware

- CONCURRENCY SEMANTICS



The “happened-before”  
relation between ops

Requires explicit  
tagging of operations

Solves the problem

# SOLUTION 1

- WHAT YOU NEED:

- CAUSAL ORDER



Off-the-shelf middleware

- CONCURRENCY SEMANTICS



The “happened-before”  
relation between ops

Requires explicit  
tagging of operations

Solves the problem

But

# SOLUTION 1

- WHAT YOU NEED:

- CAUSAL ORDER



Off-the-shelf middleware

Already tags for  
causal delivery

- CONCURRENCY SEMANTICS



The “happened-before”  
relation between ops

Requires explicit  
tagging of operations

Solves the problem

But

# SOLUTION 1

- WHAT YOU NEED:

- CAUSAL ORDER



Off-the-shelf middleware

Already tags for causal delivery

- CONCURRENCY SEMANTICS



The “happened-before” relation between ops

Tags for concurrency semantics

Requires explicit tagging of operations

Solves the problem

But

# SOLUTION 1

- WHAT YOU NEED:

- CAUSAL ORDER



Off-the-shelf middleware

Already tags for causal delivery

- CONCURRENCY SEMANTICS



The “happened-before” relation between ops

Requires explicit tagging of operations

Tags for concurrency semantics

Solves the problem

But

Duplicates effort  
More meta-data

# SOLUTION 2

- WHAT YOU NEED:

- CAUSAL ORDER



Off-the-shelf middleware

- CONCURRENCY SEMANTICS



The “happened-before”  
relation between ops

# SOLUTION 2

- WHAT YOU NEED:

- CAUSAL ORDER



Off-the-shelf middleware



Do Not  
expose

- CONCURRENCY SEMANTICS



The “happened-before”  
relation between ops



# SOLUTION 2

- WHAT YOU NEED:

- CAUSAL ORDER

- CONCURRENCY SEMANTICS



Off-the-shelf middleware

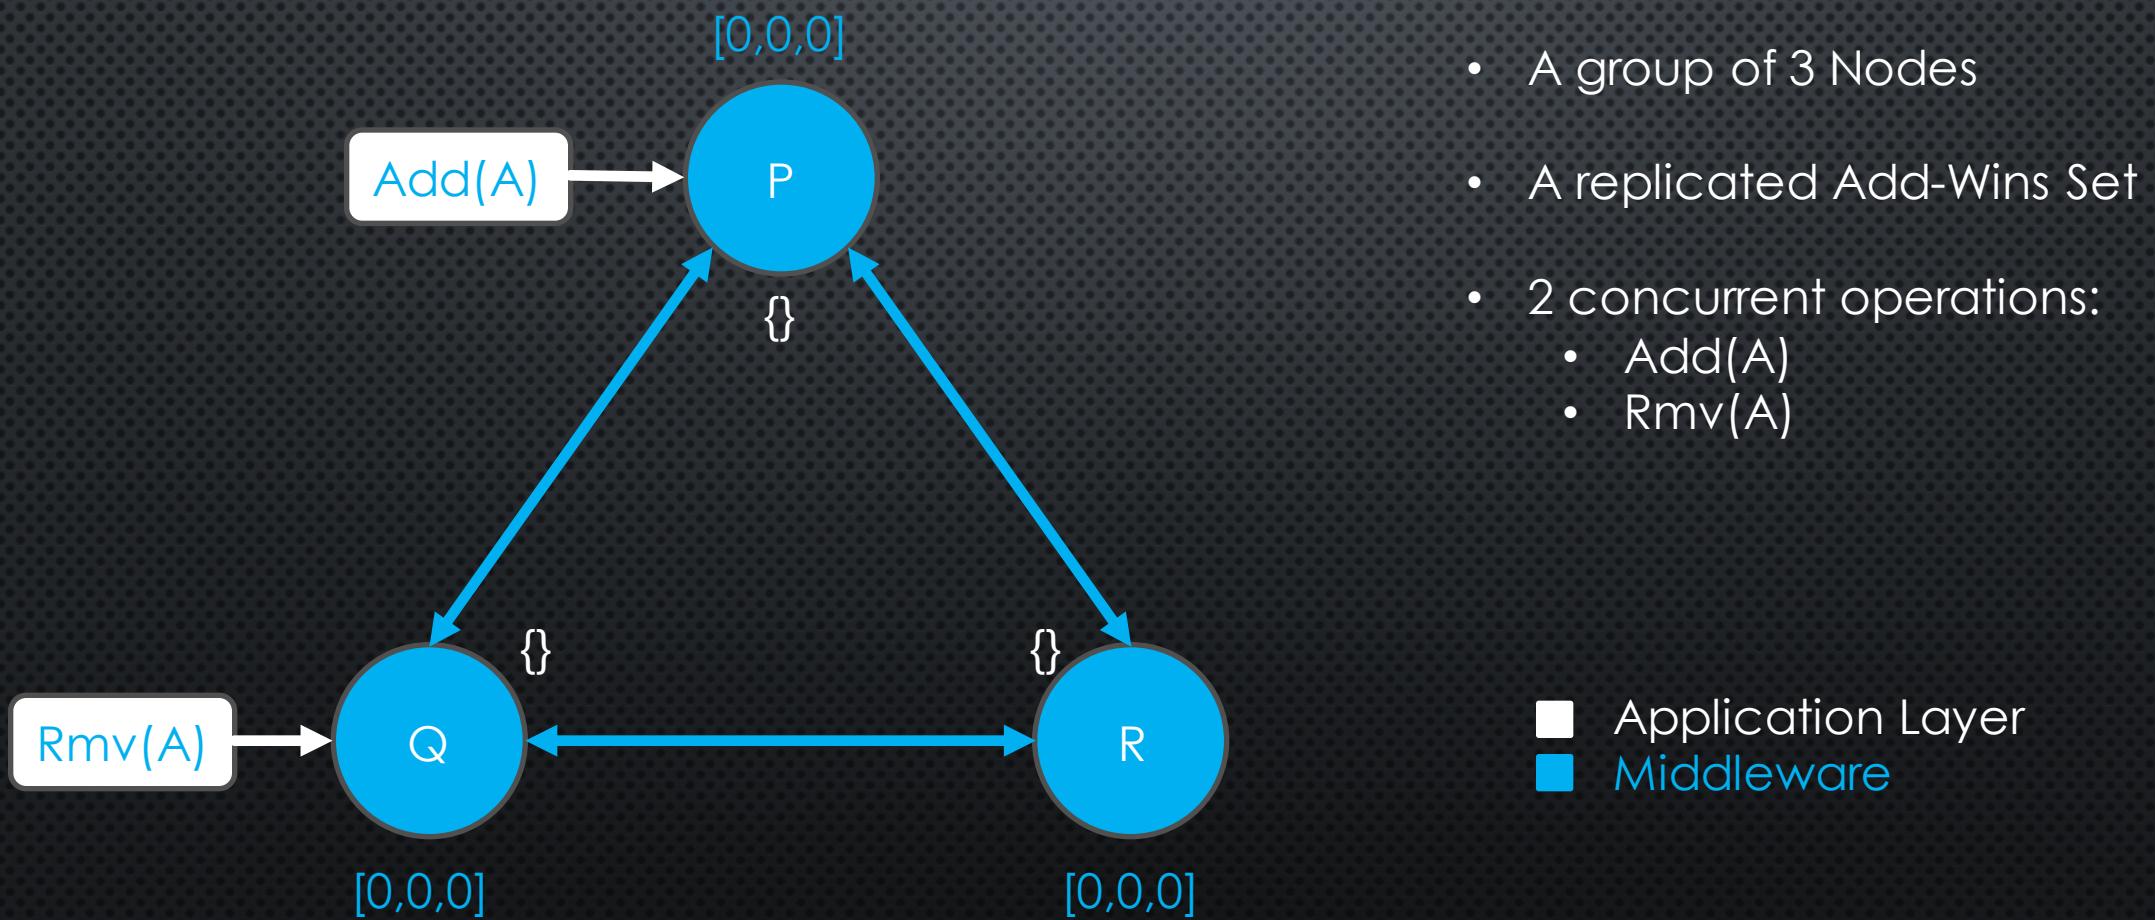
The “happened-before”  
relation between ops

Modify middleware to  
expose

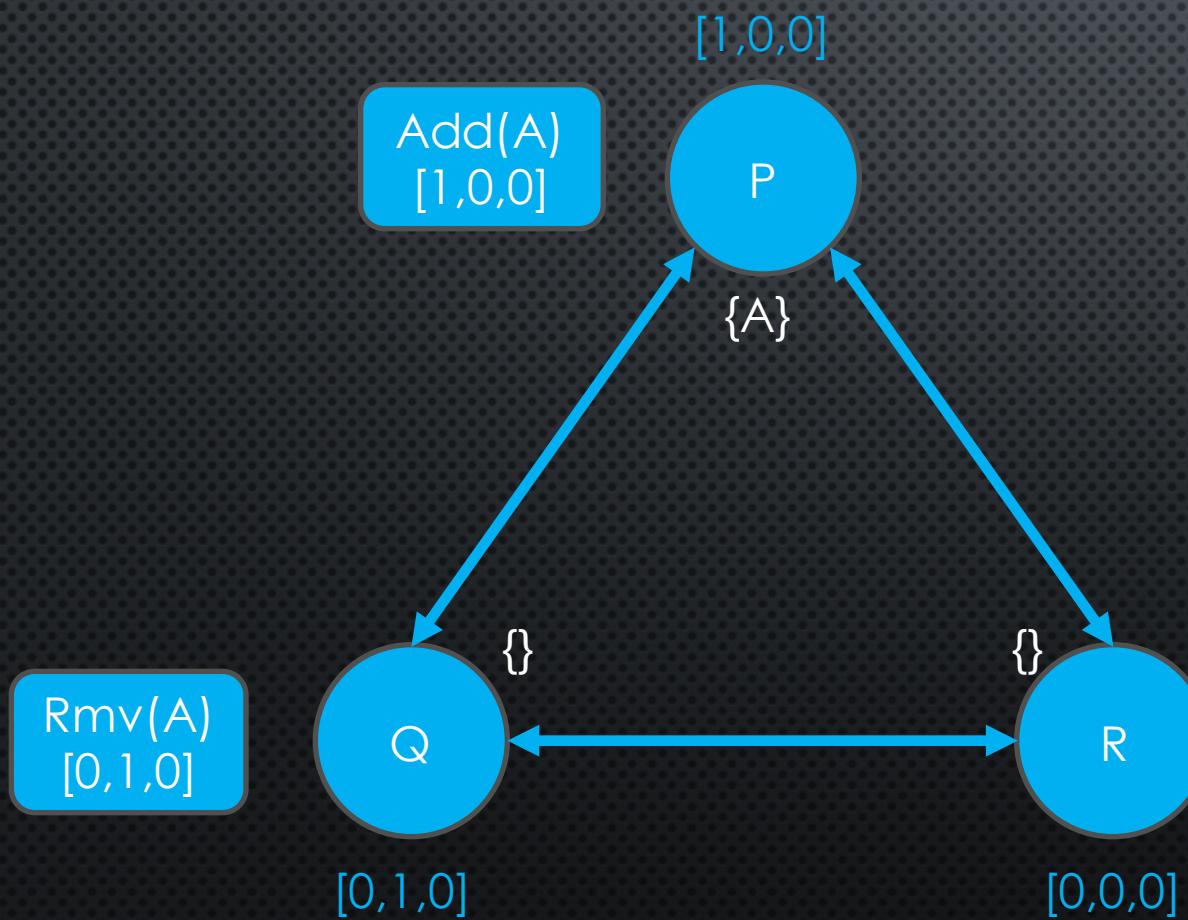
Do Not  
expose



## SOLUTION 2



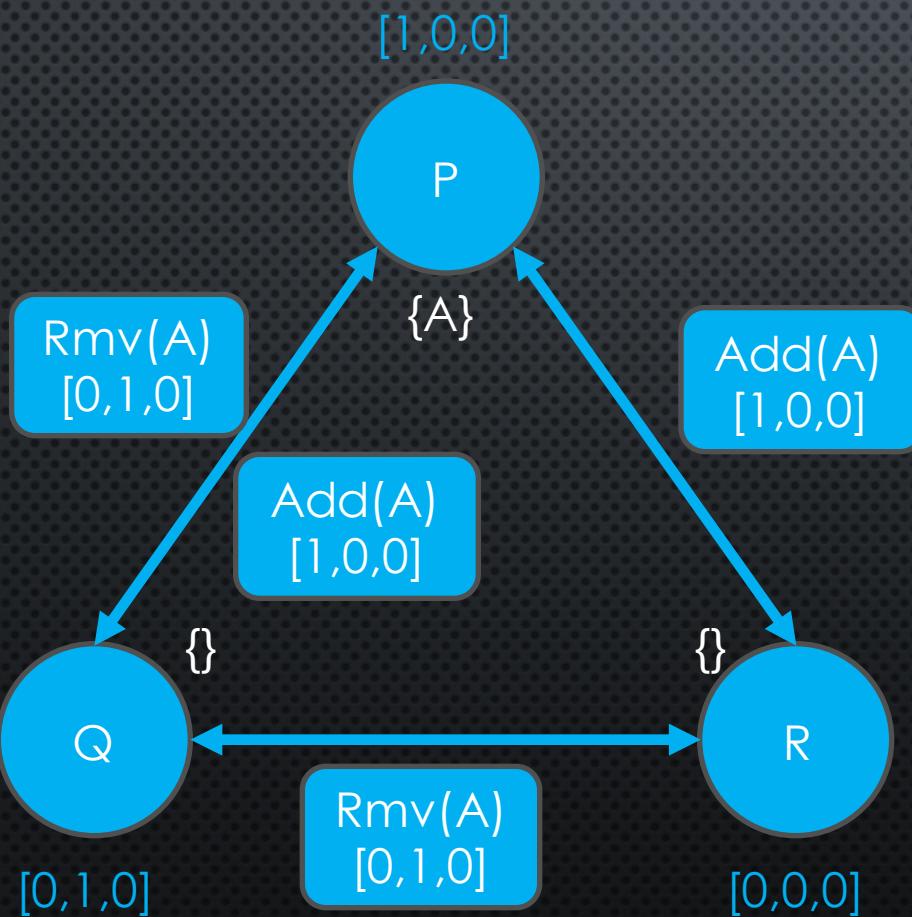
## SOLUTION 2



Scenario 1:

- P tags the op Add(A)
- Q tags the op Rmv(A)

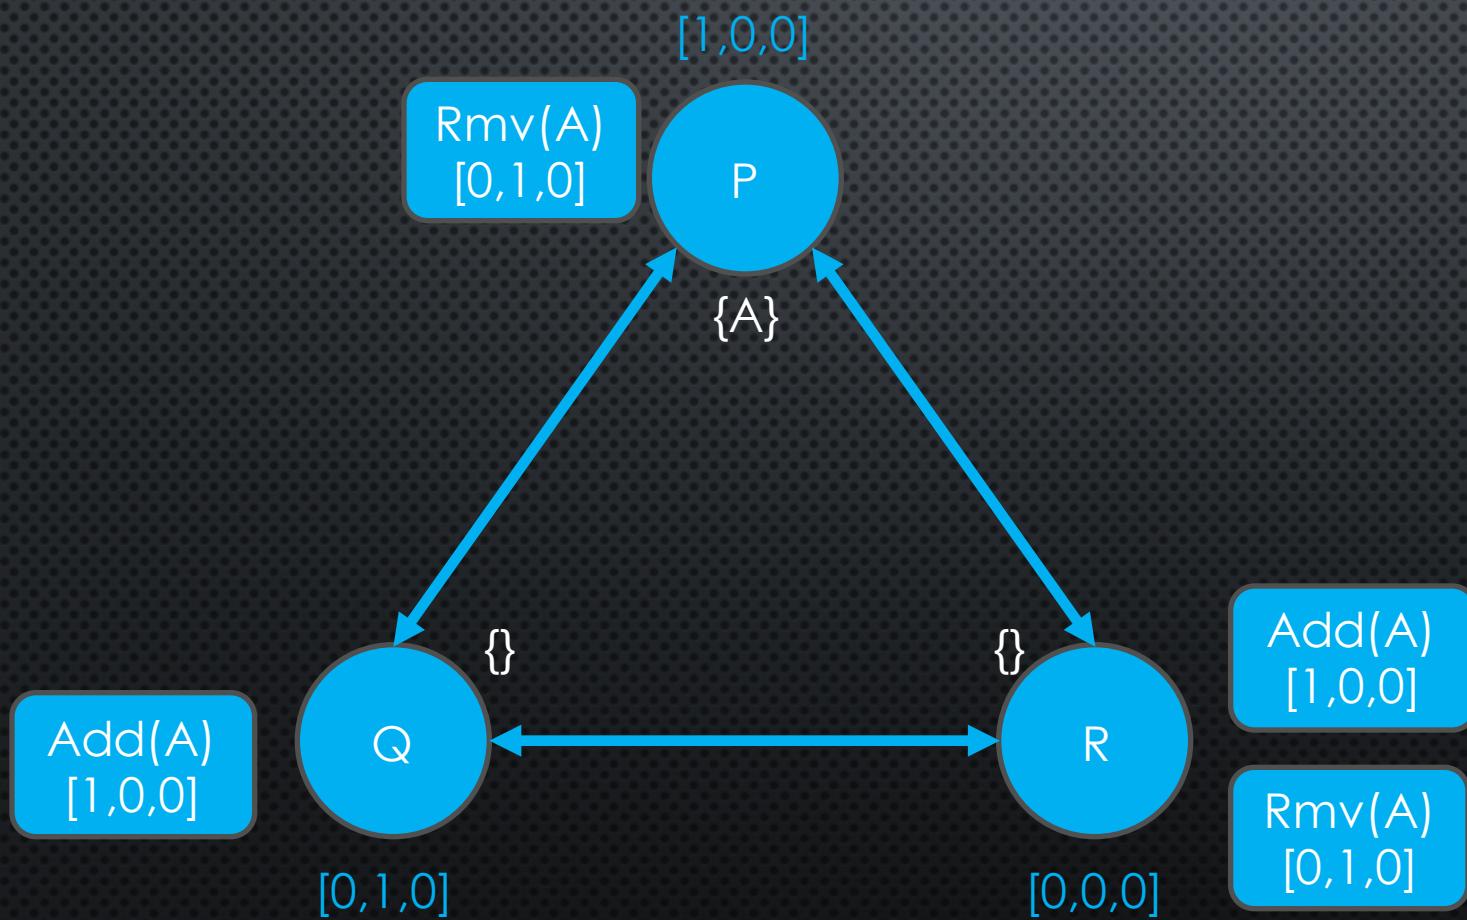
## SOLUTION 2



Scenario 1:

- P tags the op Add(A)
- Q tags the op Rmv(A)
- P and Q bcast their tagged ops

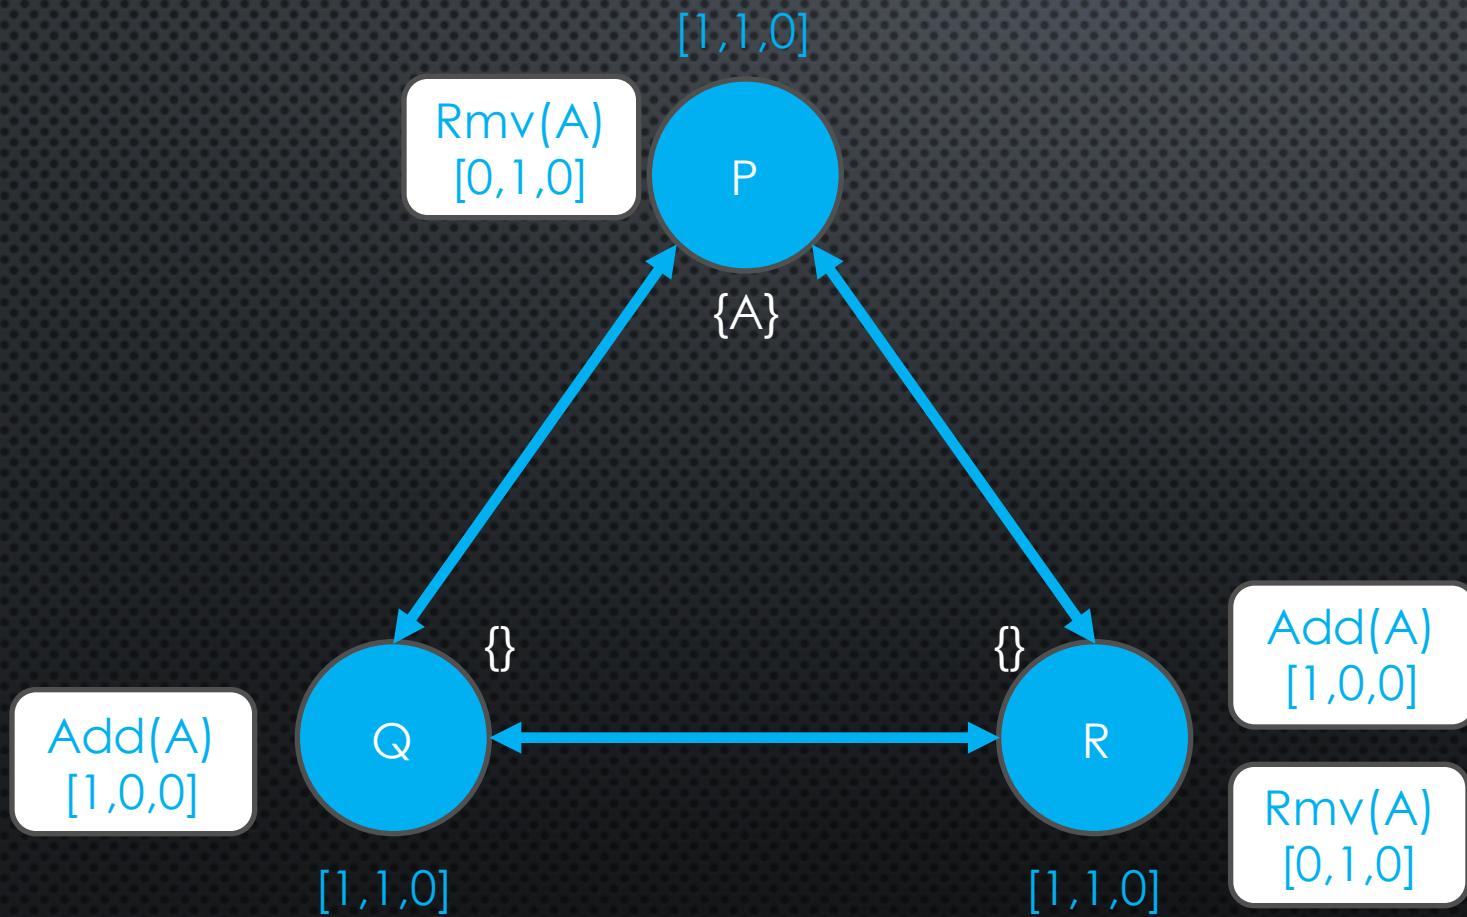
## SOLUTION 2



Scenario 1:

- P tags the op Add(A)
- Q tags the op Rmv(A)
- P and Q bcast their tagged ops

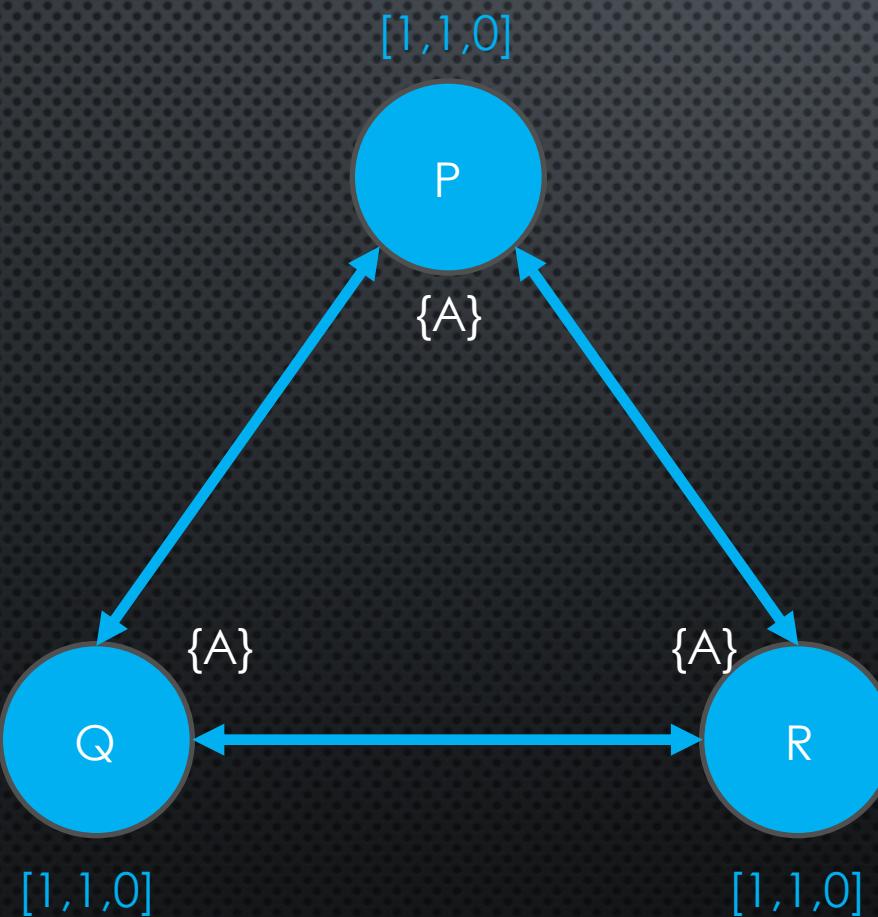
## SOLUTION 2



Scenario 1:

- P tags the op  $\text{Add}(A)$
- Q tags the op  $\text{Rmv}(A)$
- P and Q bcast their tagged ops
- All nodes deliver both ops

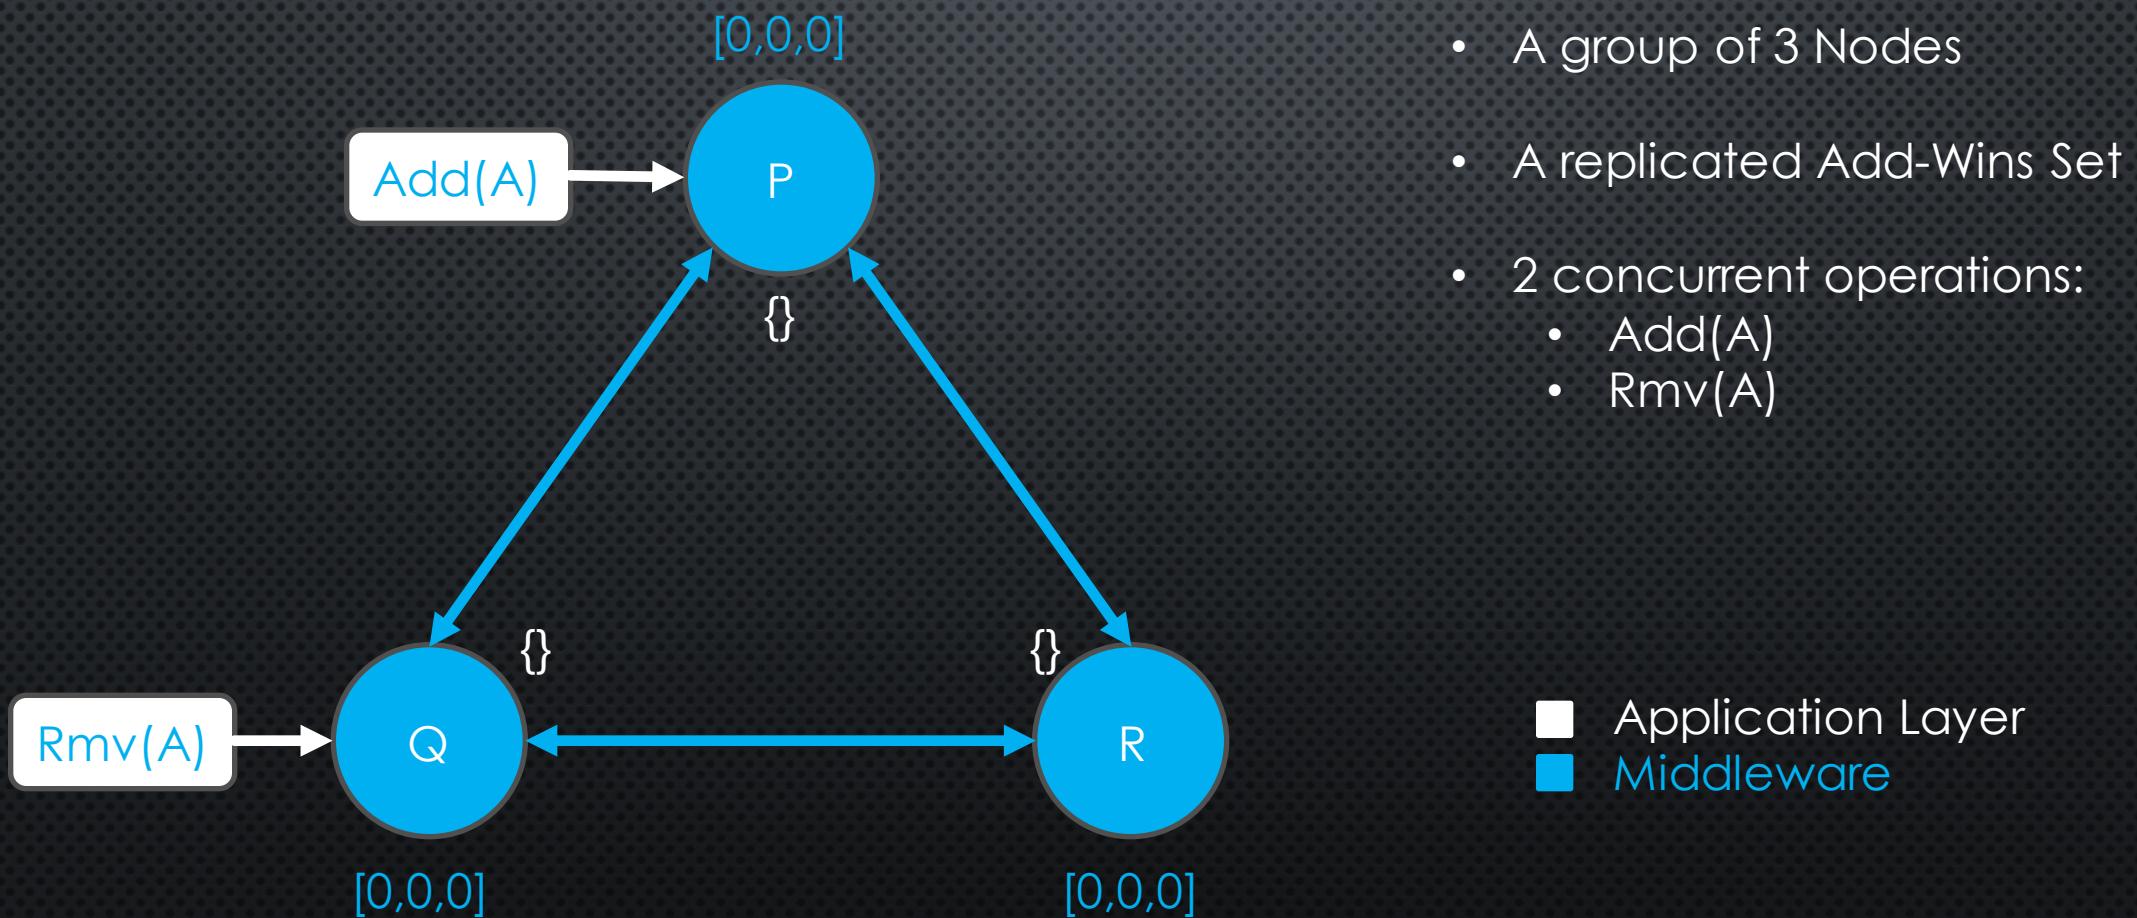
## SOLUTION 2



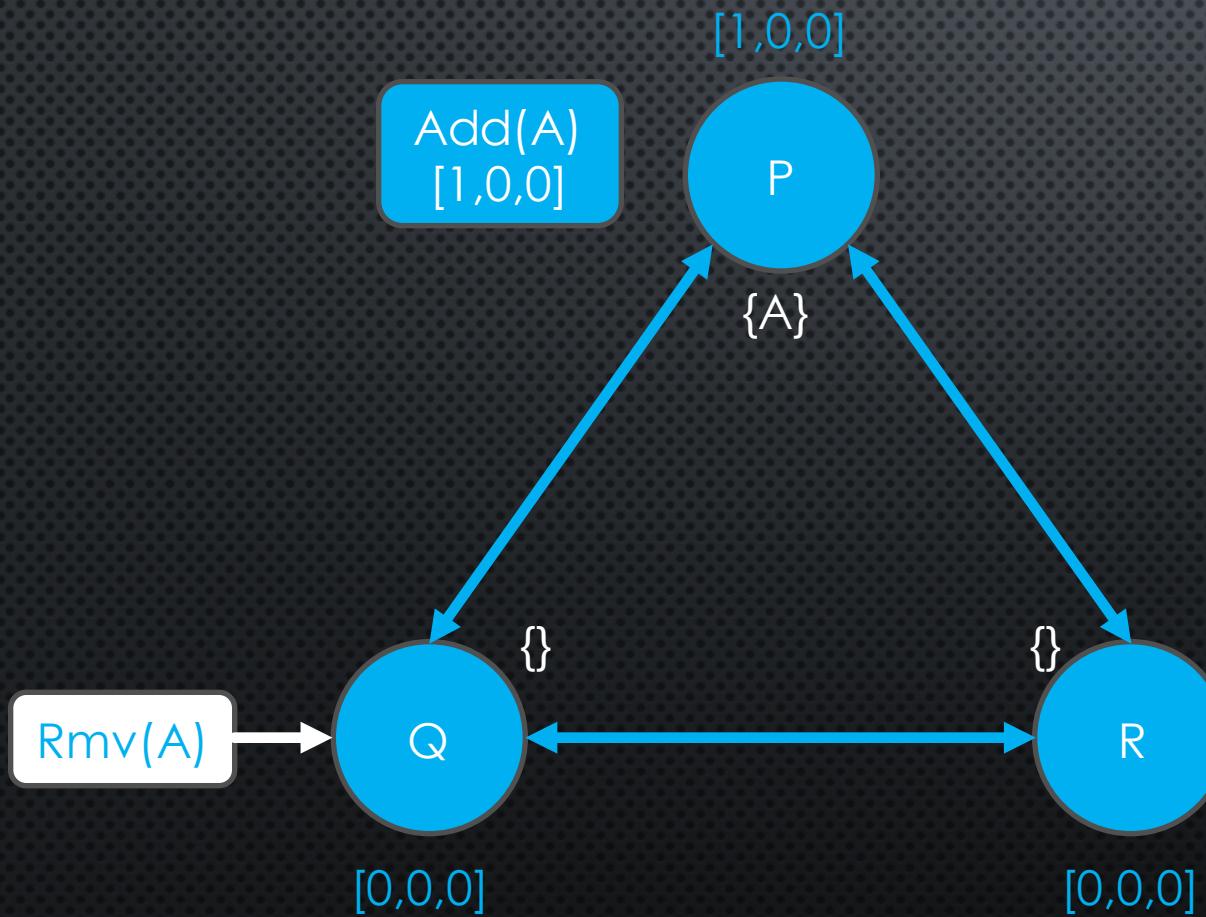
Scenario 1:

- P tags the op Add(A)
- Q tags the op Rmv(A)
- P and Q bcast their tagged ops
- All nodes deliver both ops
- [1,0,0] and [0,1,0] are concurrent
  - Add-Wins Semantics
  - A is there

## SOLUTION 2



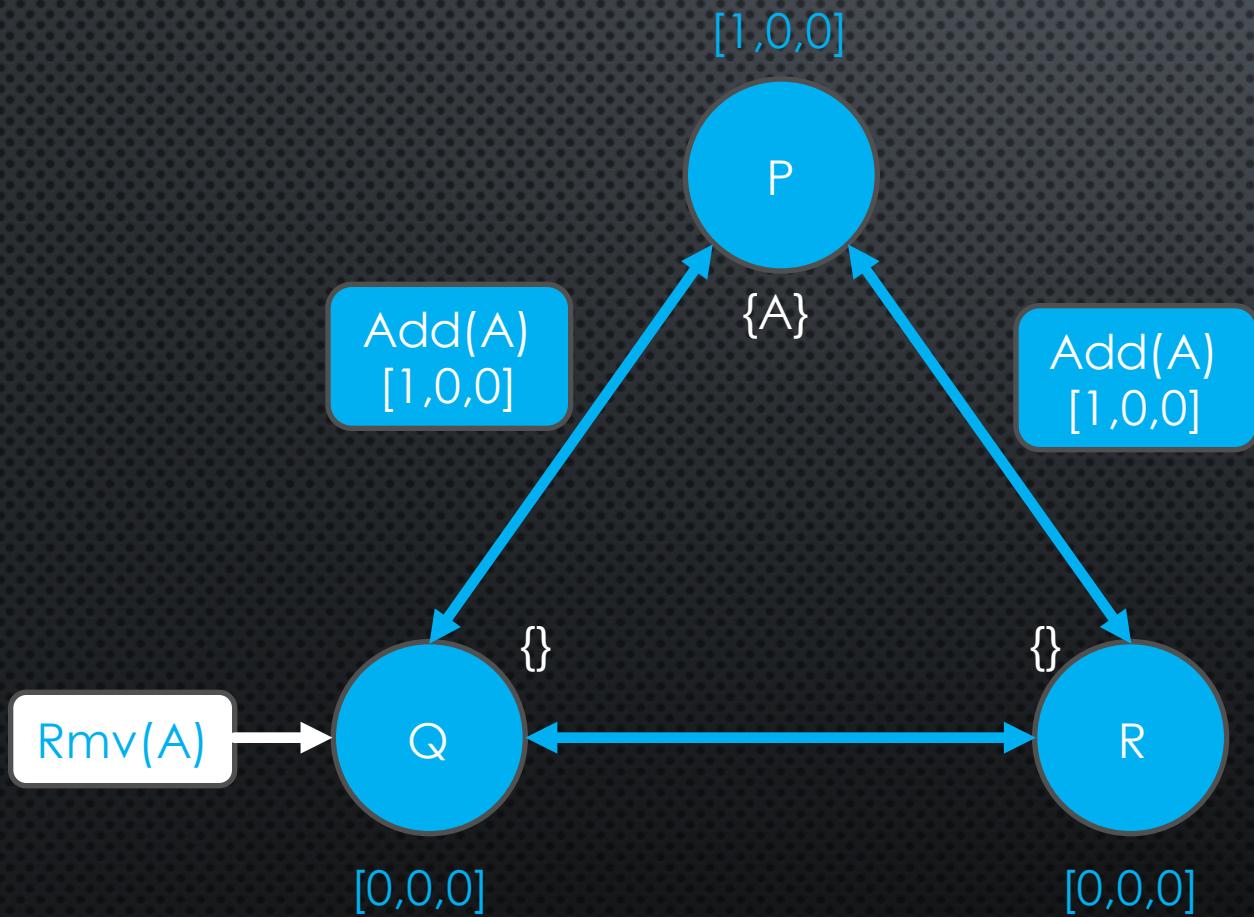
## SOLUTION 2



Scenario 2:

- P tags the op Add(A)
- Rmv(A) is still in the queue waiting to be processed by Q

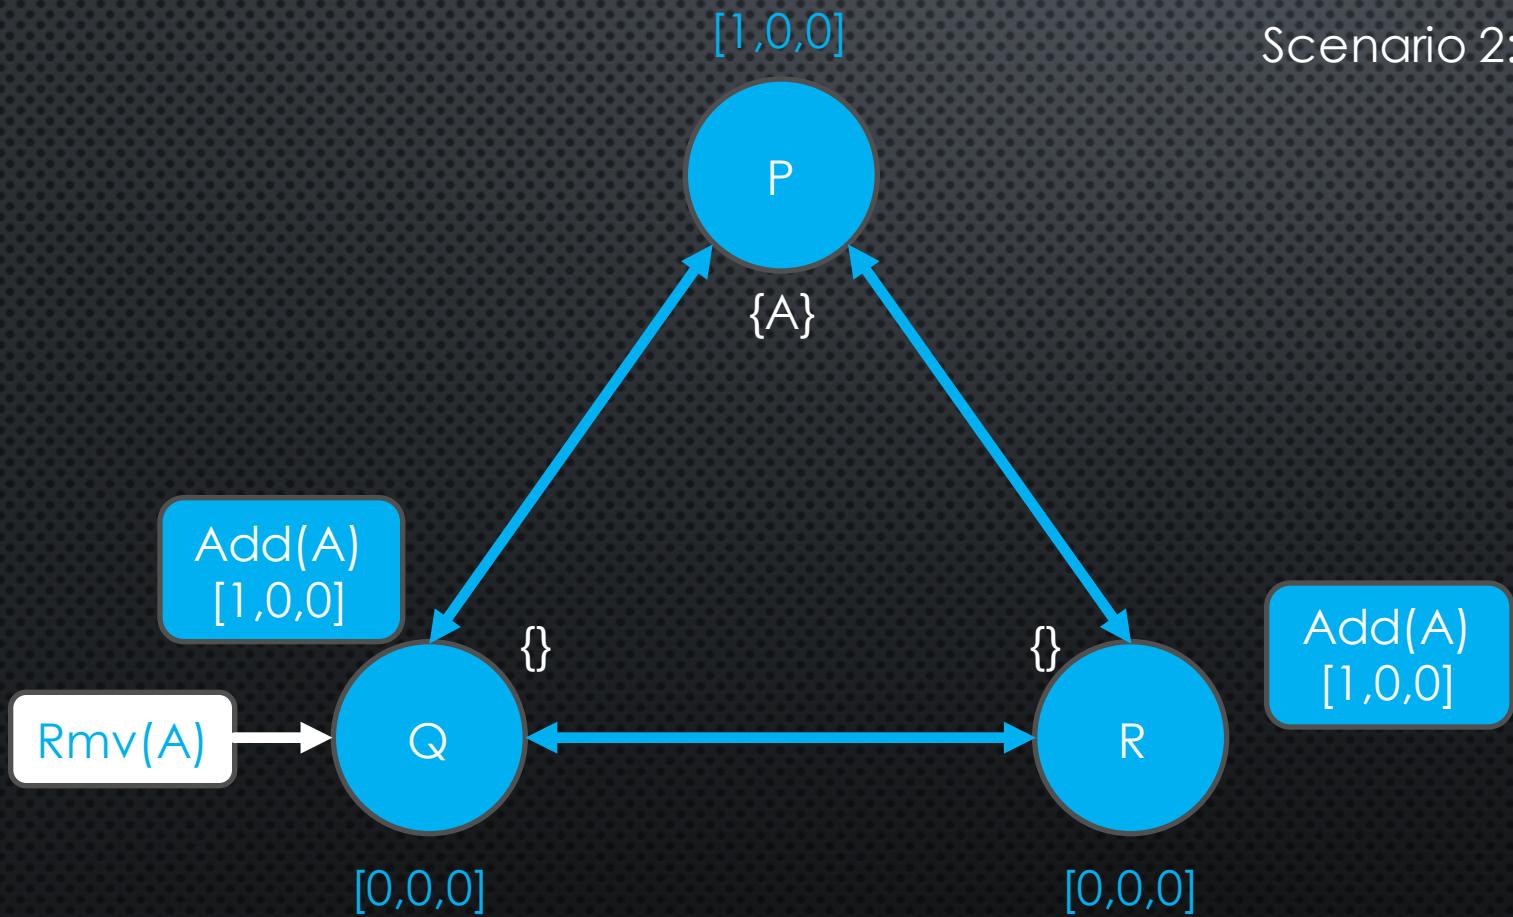
## SOLUTION 2



Scenario 2:

- P tags the op **Add(A)**
- **Rmv(A)** is still in the queue waiting to be processed by Q
- P broadcasts their tagged op

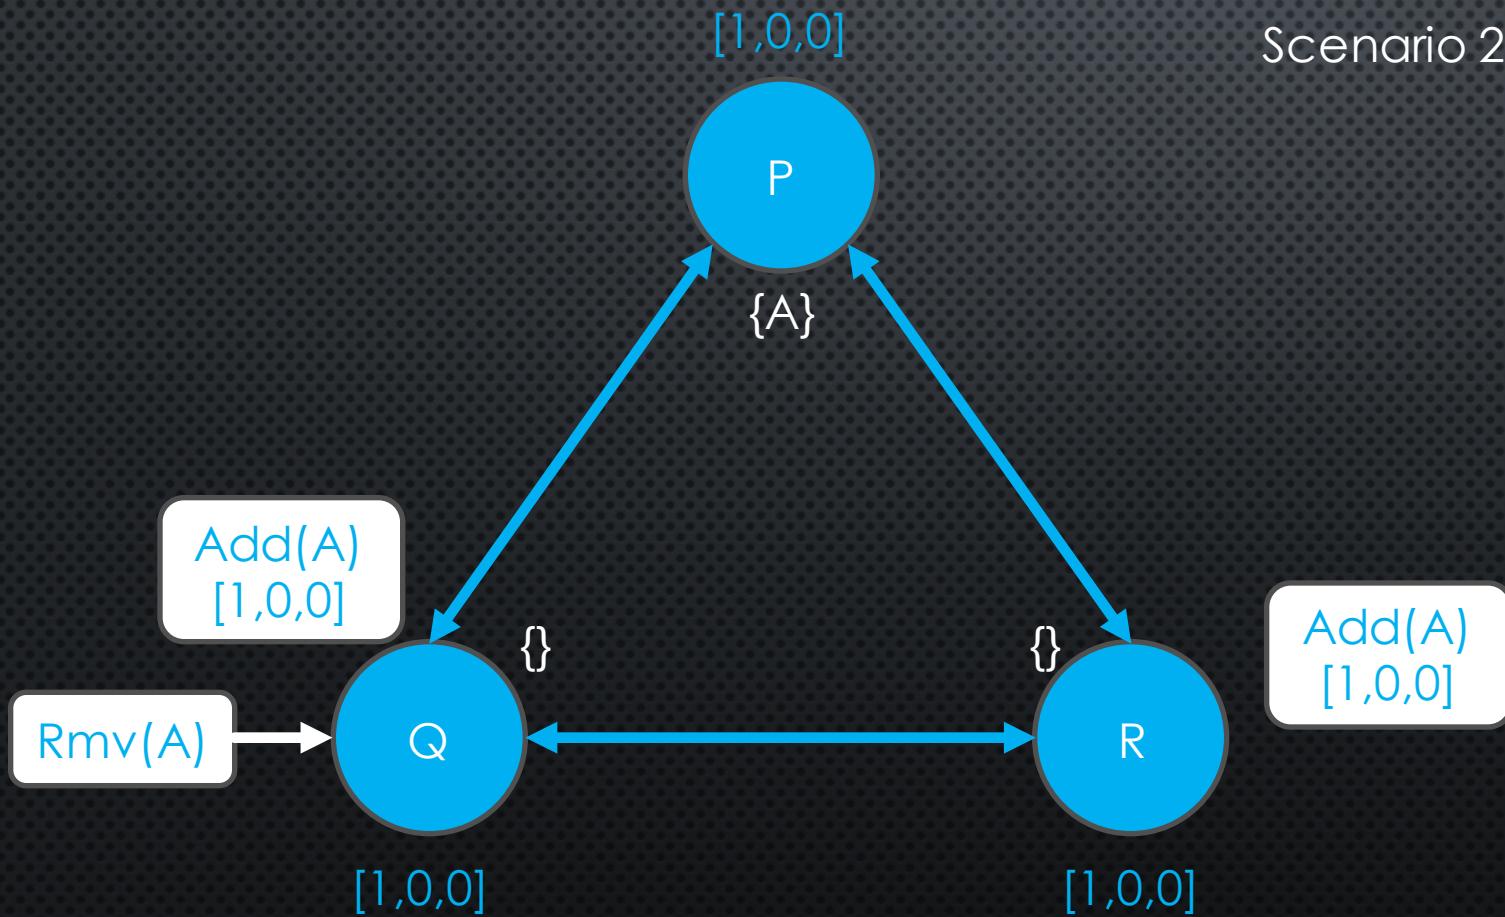
## SOLUTION 2



Scenario 2:

- P tags the op Add(A)
- Rmv(A) is still in the queue waiting to be processed by Q
- P broadcasts their tagged op

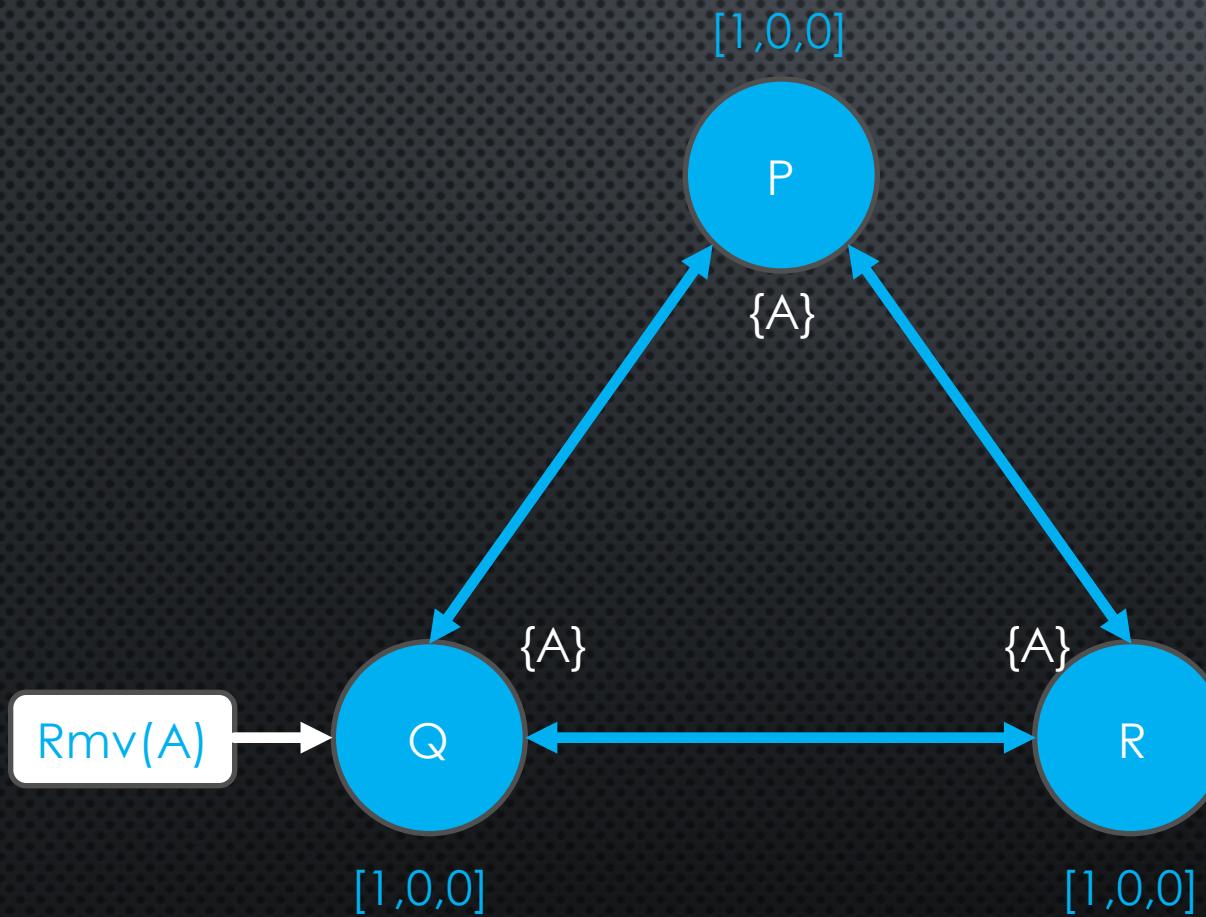
## SOLUTION 2



Scenario 2:

- P tags the op Add(A)
- Rmv(A) is still in the queue waiting to be processed by Q
- P broadcasts their tagged op
- All nodes deliver Add(A)

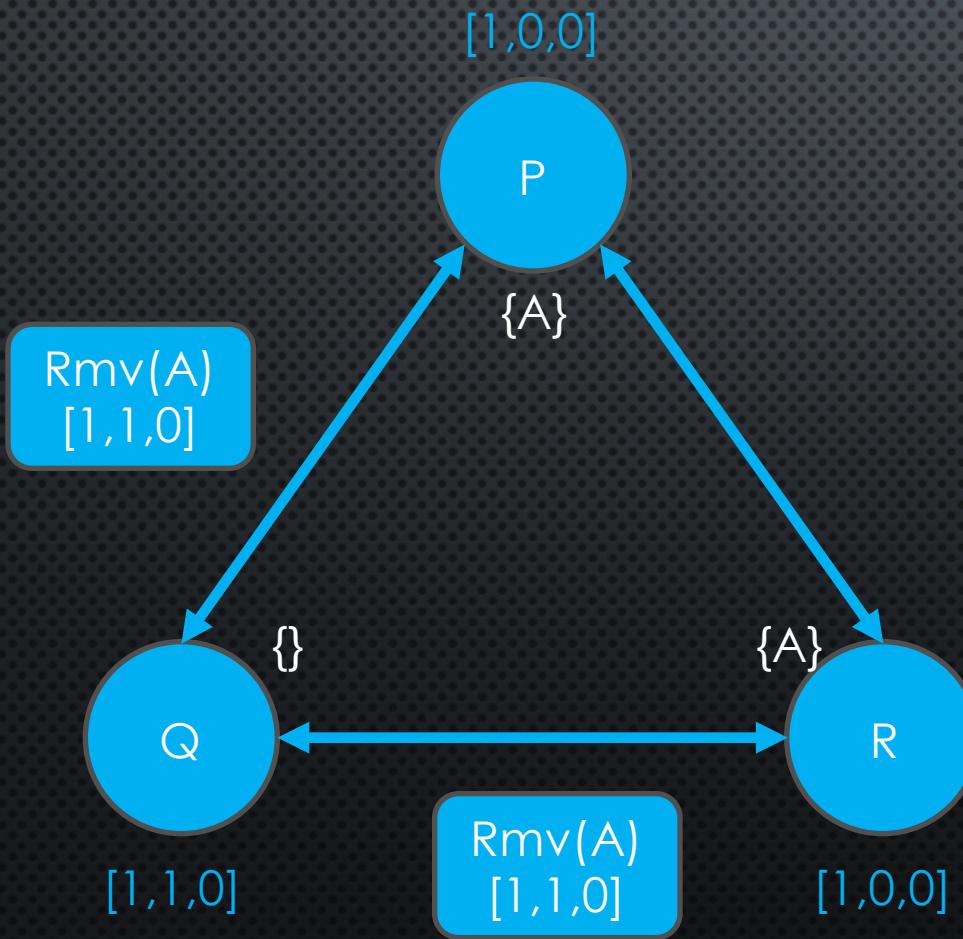
## SOLUTION 2



Scenario 2:

- P tags the op  $\text{Add}(A)$
- $\text{Rmv}(A)$  is still in the queue waiting to be processed by Q
- P broadcasts their tagged op
- All nodes deliver  $\text{Add}(A)$

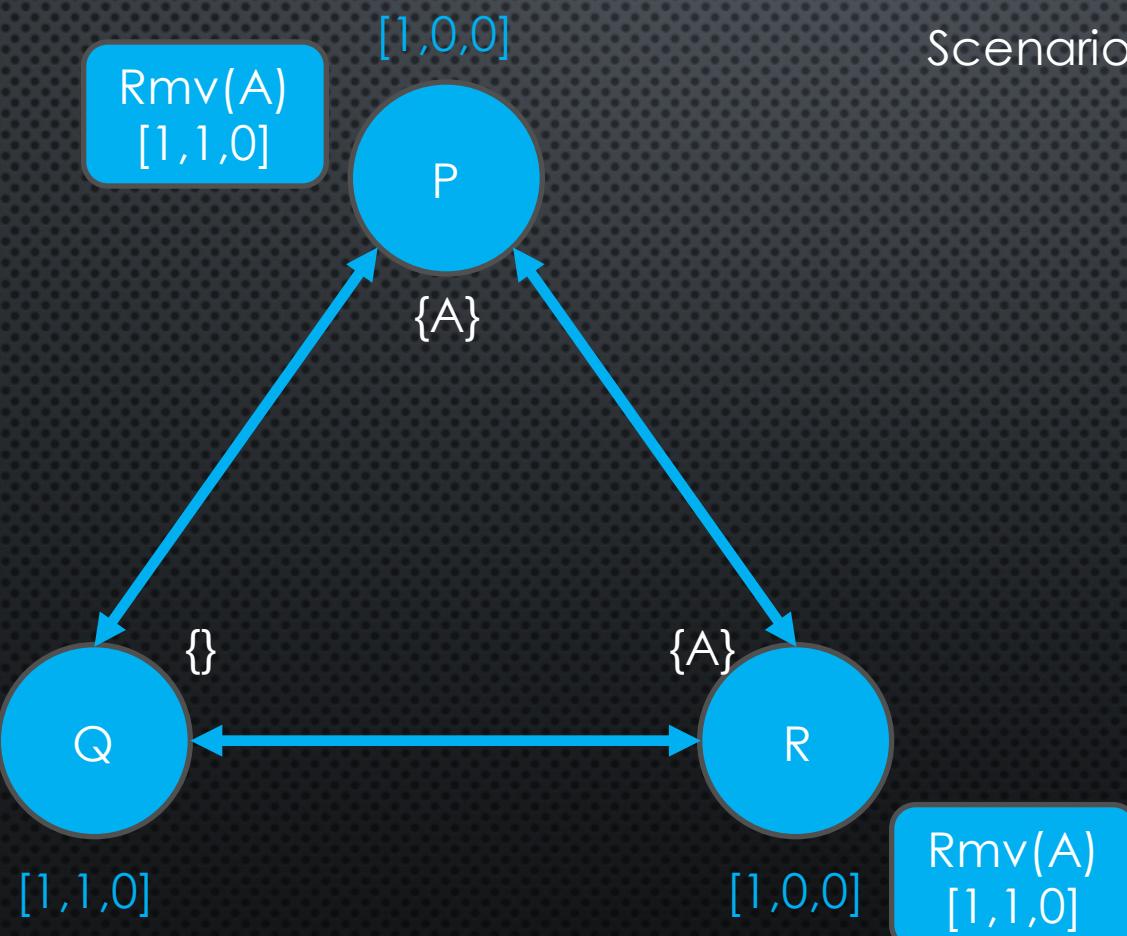
## SOLUTION 2



Scenario 2:

- P tags the op Add(A)
- Rmv(A) is still in the queue waiting to be processed by Q
- P bcasts their tagged op
- All nodes deliver Add(A)
- Q dequeues Rmv(A) , tags and bcasts it

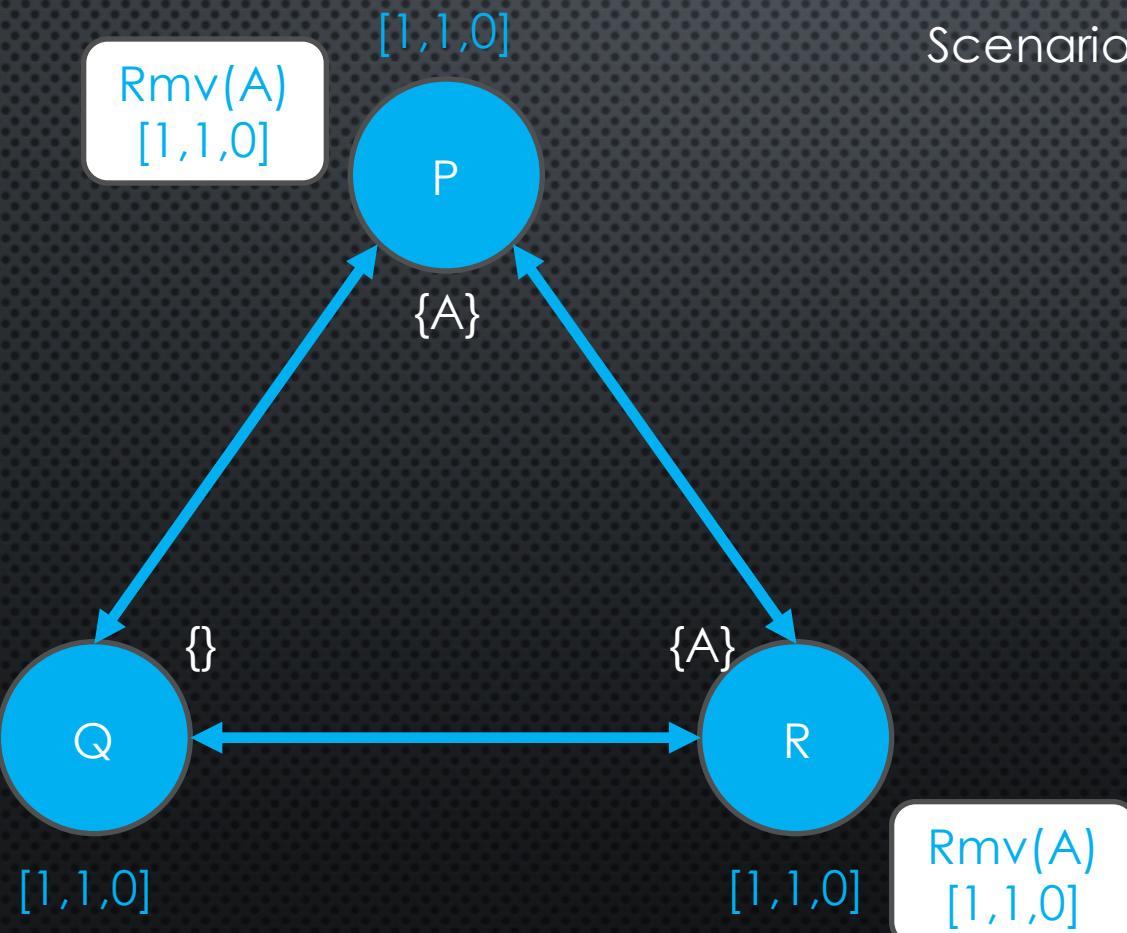
## SOLUTION 2



Scenario 2:

- P tags the op  $\text{Add}(A)$
- $\text{Rmv}(A)$  is still in the queue waiting to be processed by Q
- P bcasts their tagged op
- All nodes deliver  $\text{Add}(A)$
- Q dequeues  $\text{Rmv}(A)$ , tags and bcasts it

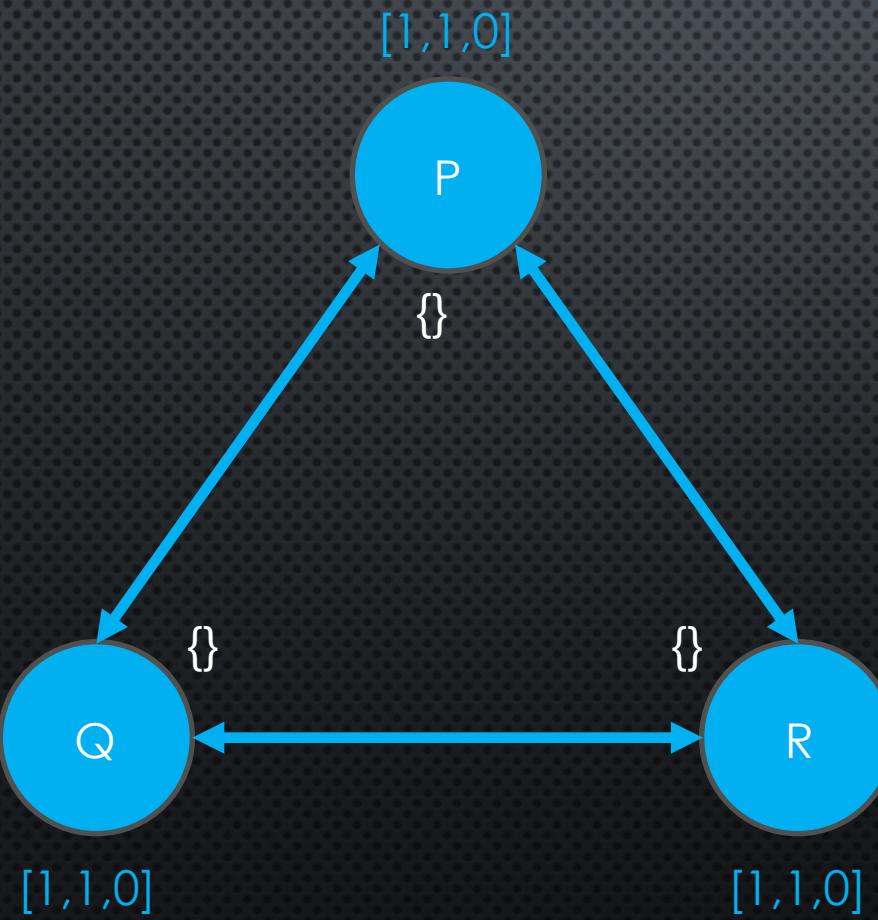
## SOLUTION 2



Scenario 2:

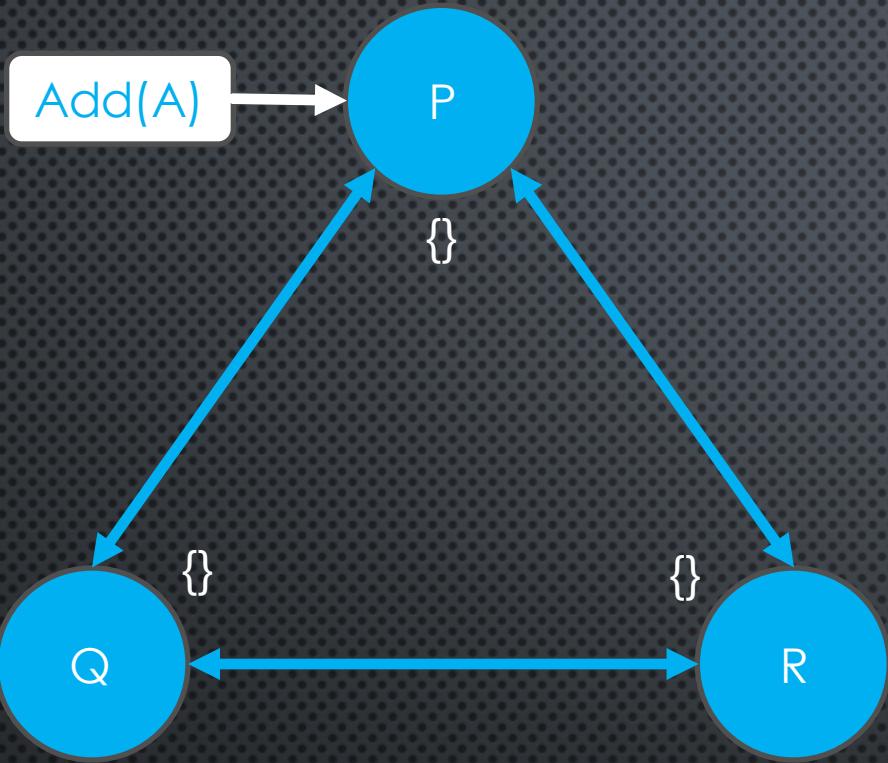
- P tags the op Add(A)
- Rmv(A) is still in the queue waiting to be processed by Q
- P bcasts their tagged op
- All nodes deliver Add(A)
- Q dequeues Rmv(A) , tags and bcasts it
- Rmv(A) delivered at P and R

## SOLUTION 2



Scenario 2:

- P tags the op Add(A)
- Rmv(A) is still in the queue waiting to be processed by Q
- P bcasts their tagged op
- All nodes deliver Add(A)
- Q dequeues Rmv(A) , tags and bcasts it
- Rmv(A) delivered at P and R
- [1,0,0] happened-before [1,1,0]
  - A is not there (**expected {A}**)



### Scenario 1:

- Add(A) and Rmv(A) tagged as concurrent
- Final State: {A}

### Solution 2: not a solution

- Does not “care” about concurrent operations
- Orders concurrent ops based on their delivery order
  - Could order concurrent ops as one happening before the other (Scenario 2)

### Scenario 2:

- Rmv(A) tagged as in the future of Add(A)
- Final State: {}

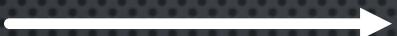
# SOLUTION 3

- BETTER CHARACTERIZATION OF THE HAPPENED-BEFORE RELATION BETWEEN OPS
  - TAG BASED ON WHAT WAS DELIVERED AT THE APPLICATION LEVEL
- AS DELIVERY HAPPENS AT THE APPLICATION LEVEL
  - TAG AT THE APPLICATION LEVEL

# SOLUTION 3

- WHAT YOU NEED:

- CAUSAL ORDER



- CONCURRENCY SEMANTICS



Tag operations  
at app level

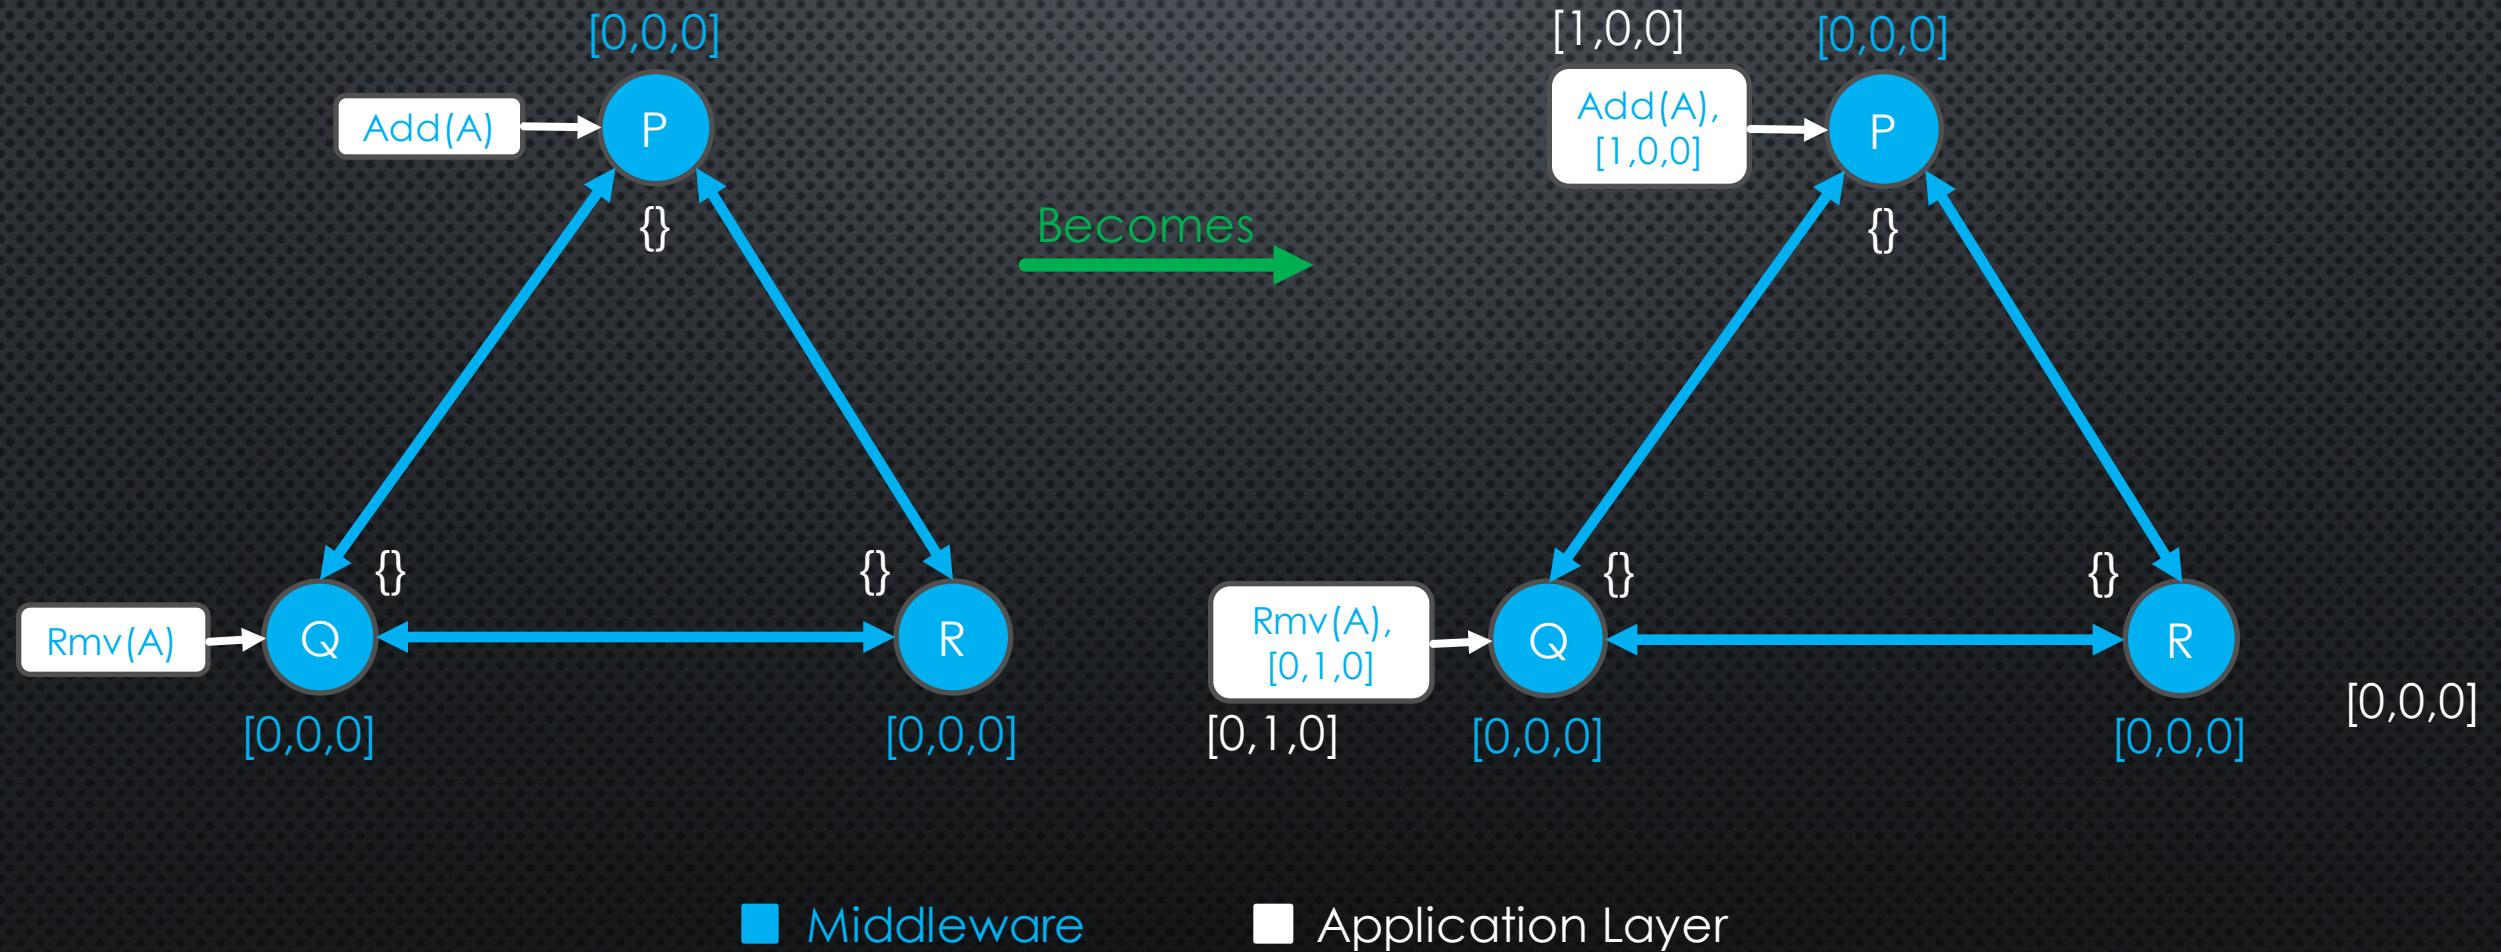
Use those tags  
for causal  
delivery

Off-the-shelf middleware

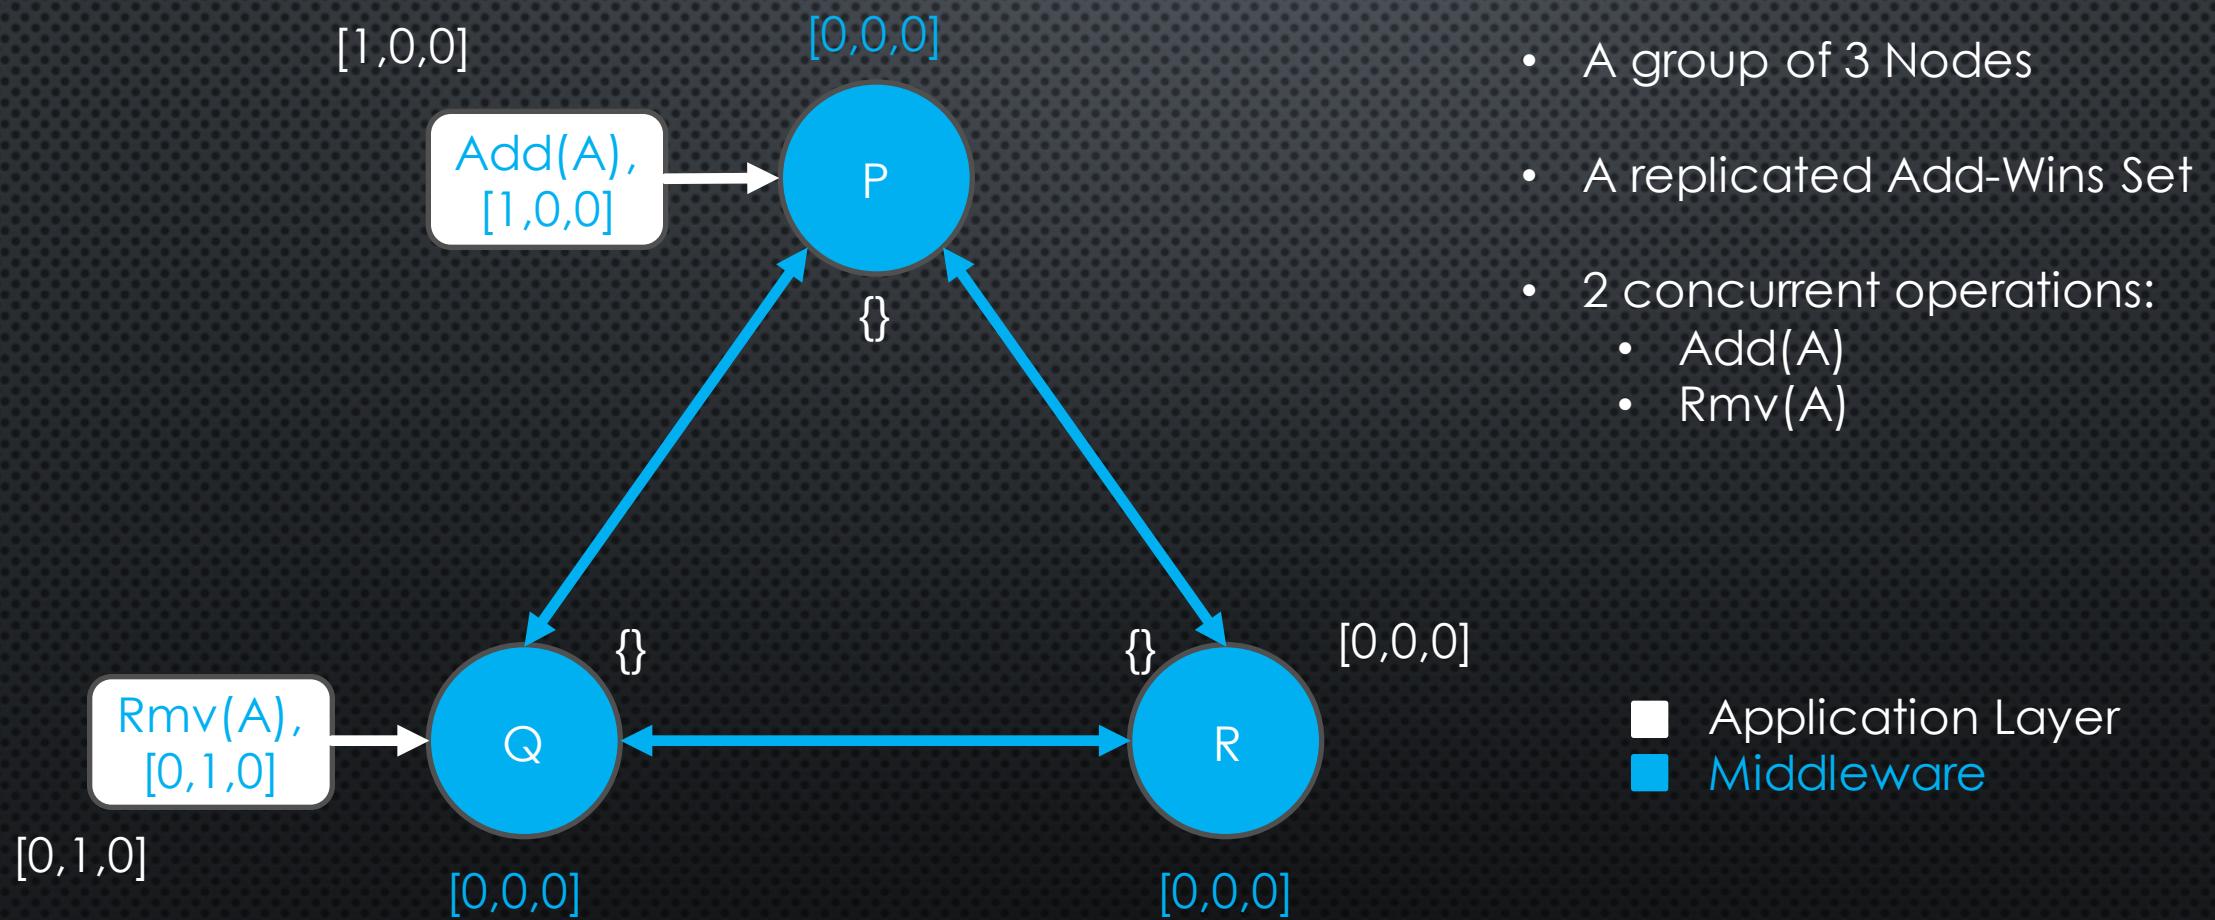
The “happened-before”  
relation between ops

Requires tagging  
operations

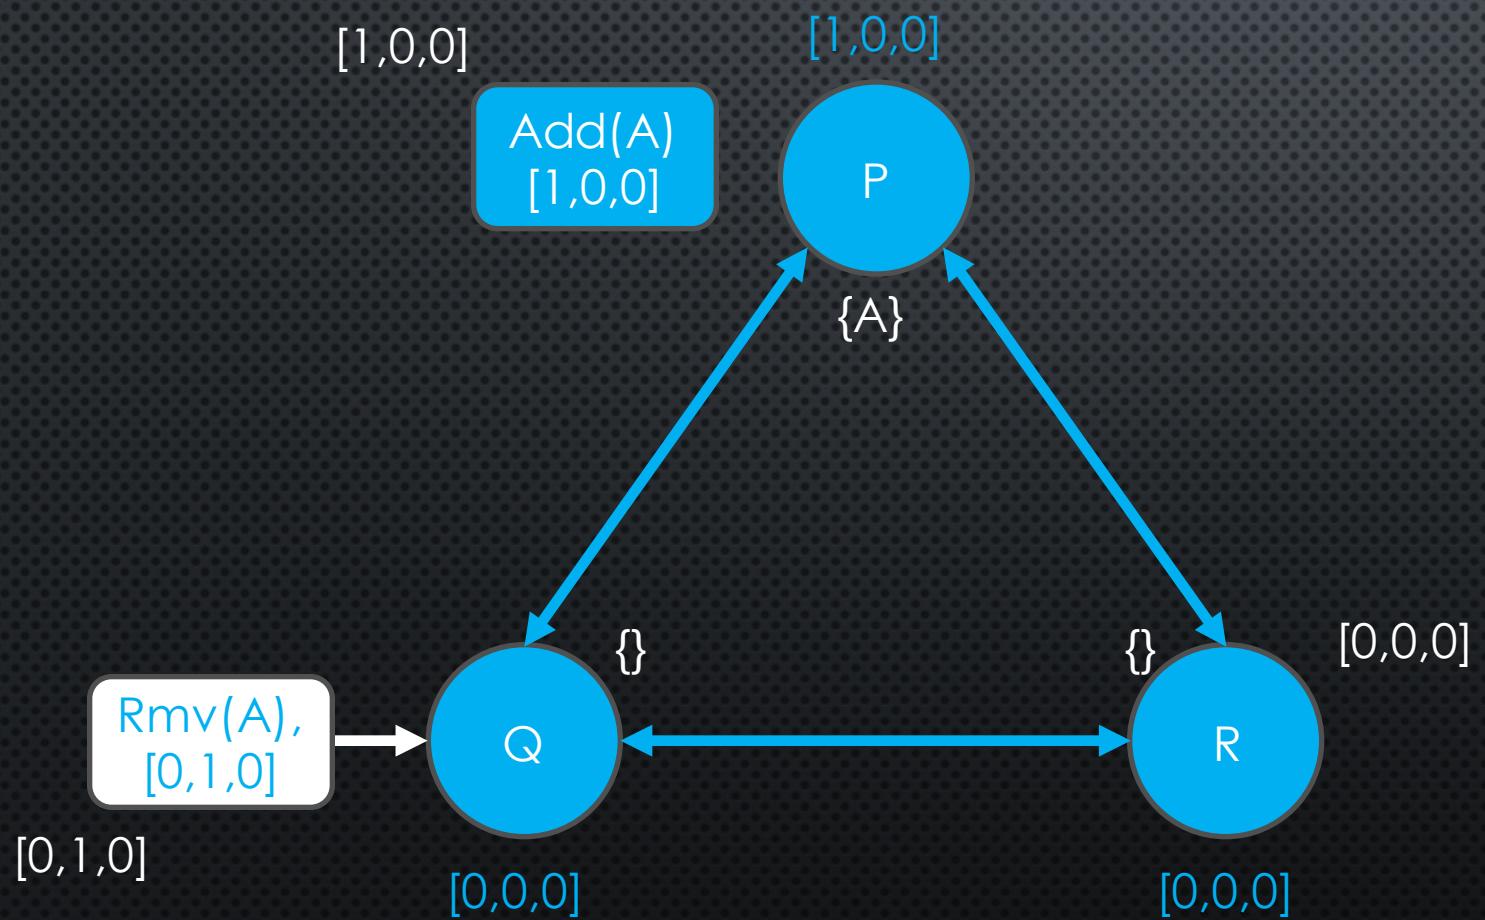
# SOLUTION 3



# SOLUTION 3



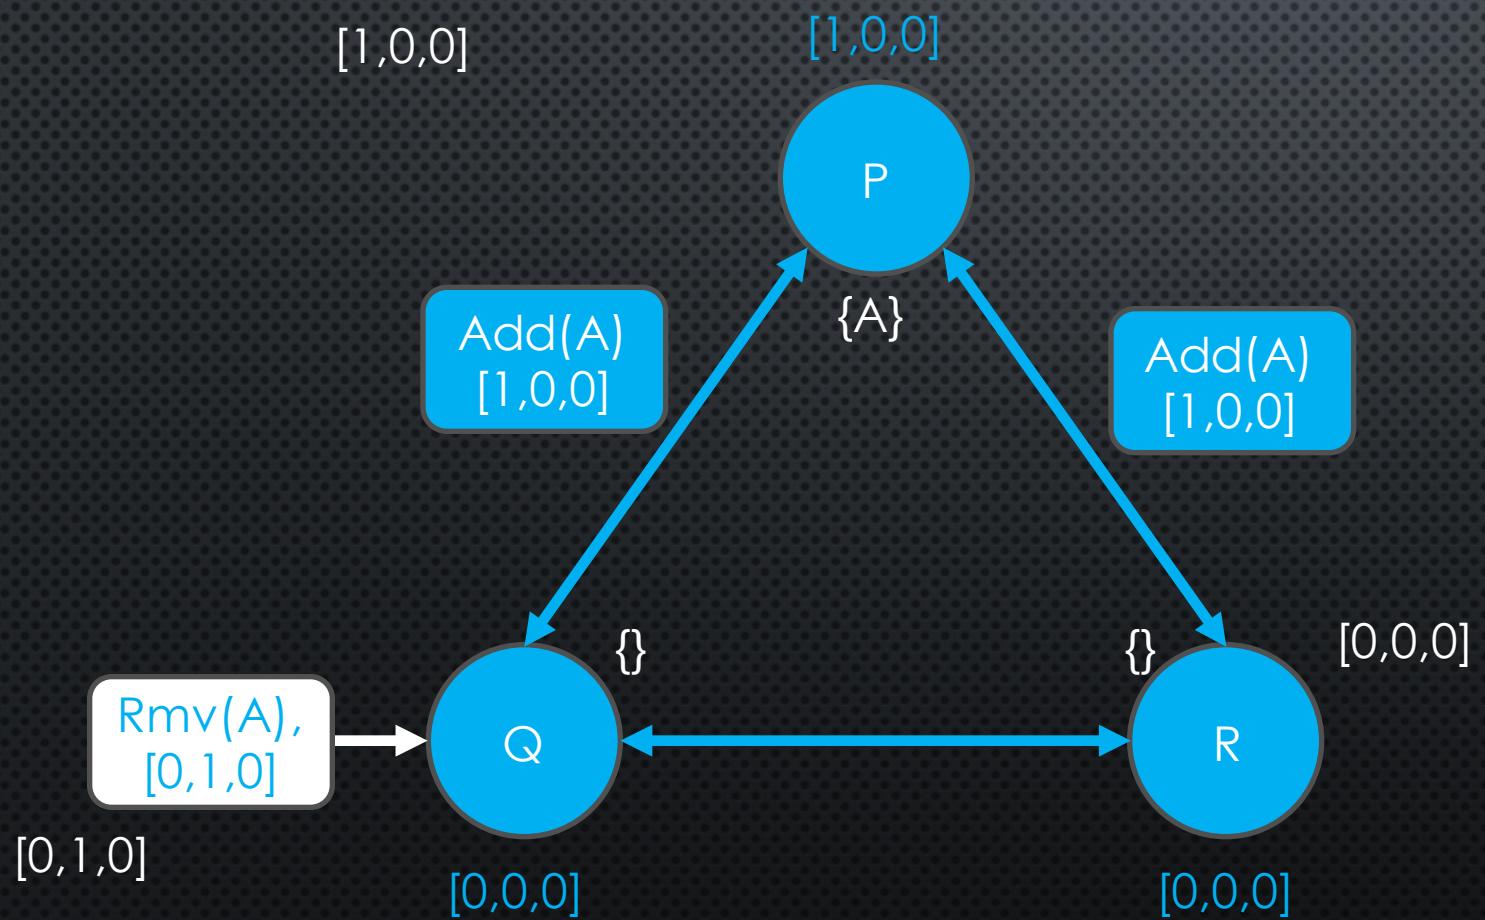
# SOLUTION 3



Scenario 2:

- Both ops tagged at App level
- P bcasts Add(A)
- Rmv(A) is still in the queue waiting to be bcast by Q

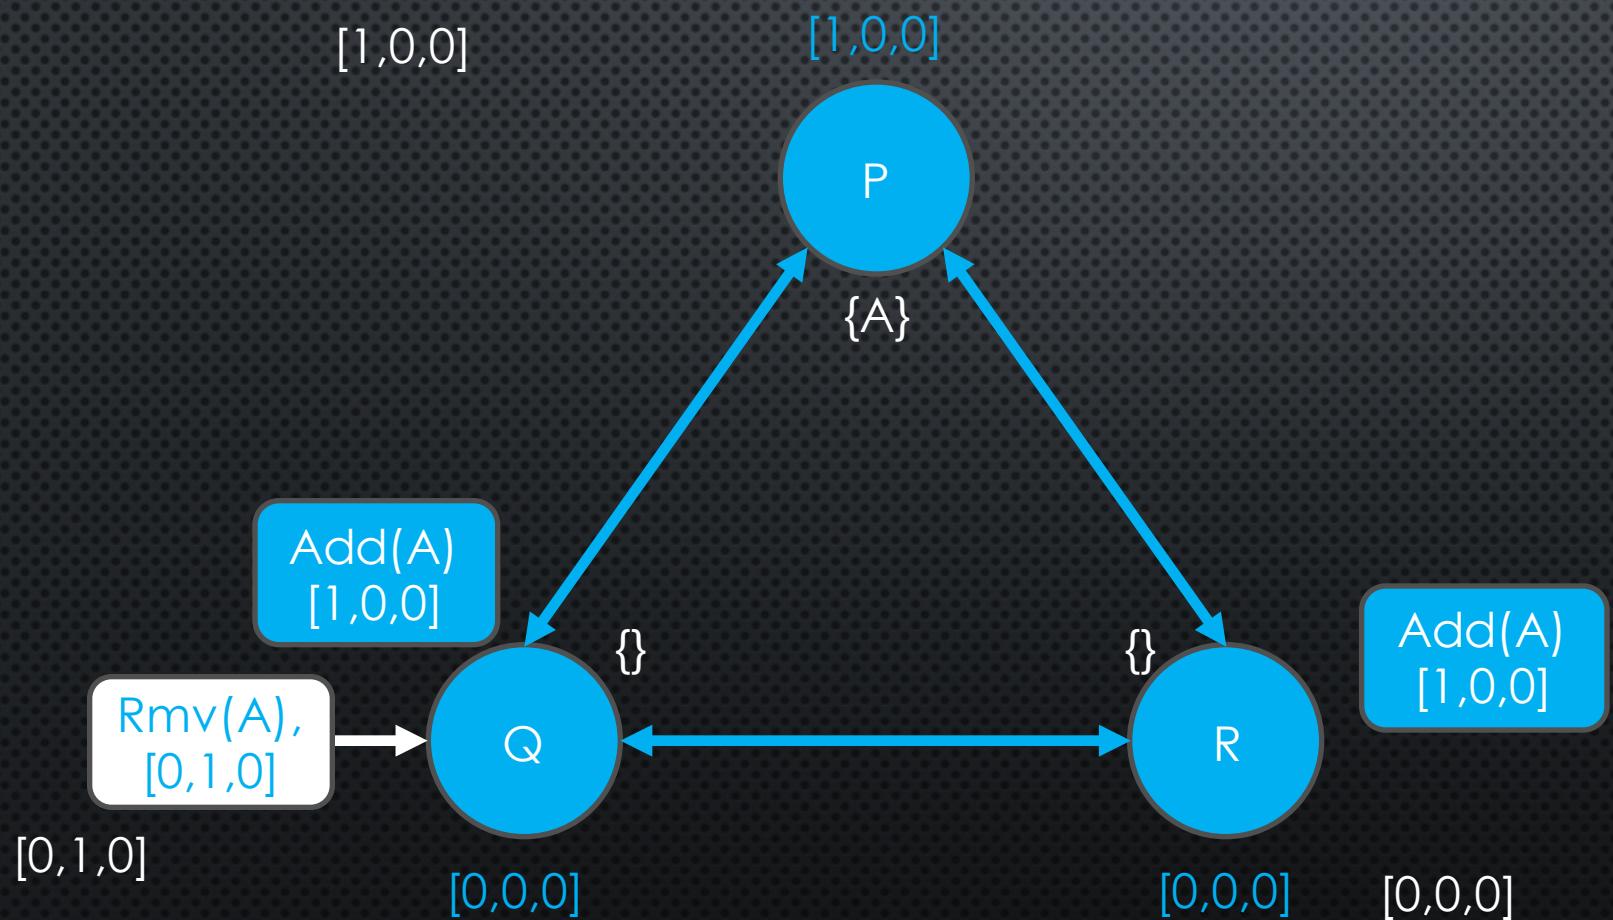
# SOLUTION 3



Scenario 2:

- Both ops tagged at App level
- P bcasts Add(A)
- Rmv(A) is still in the queue waiting to be bcast by Q

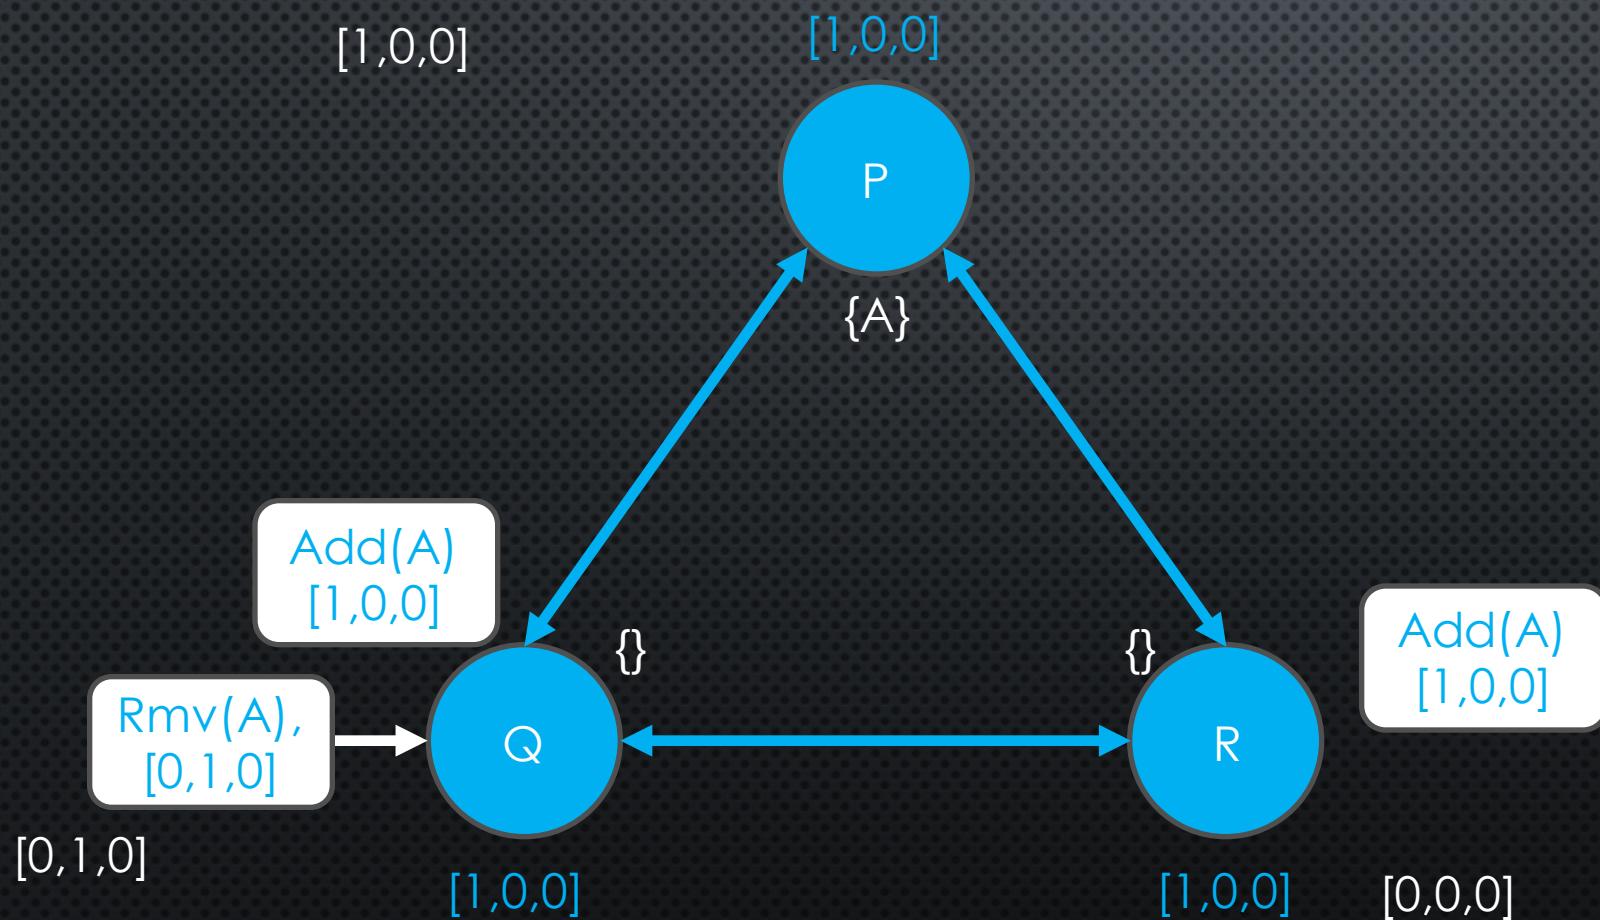
# SOLUTION 3



Scenario 2:

- Both ops tagged at App level
- P bcasts Add(A)
- Rmv(A) is still in the queue waiting to be bcast by Q

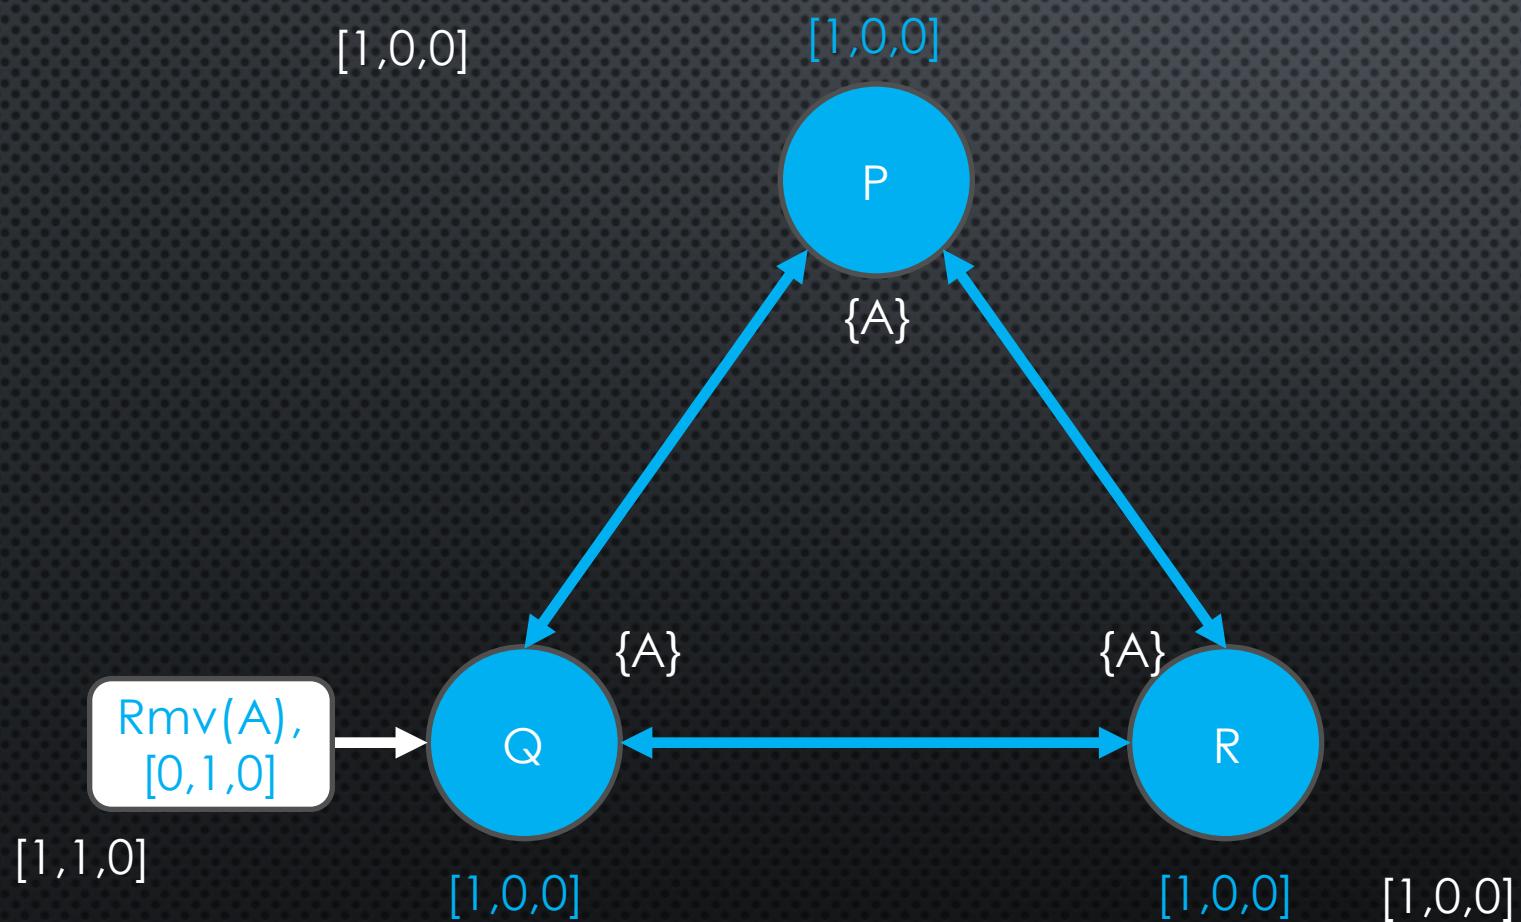
# SOLUTION 3



Scenario 2:

- Both ops tagged at App level
- P bcasts Add(A)
- Rmv(A) is still in the queue waiting to be bcast by Q
- All nodes deliver Add(A)

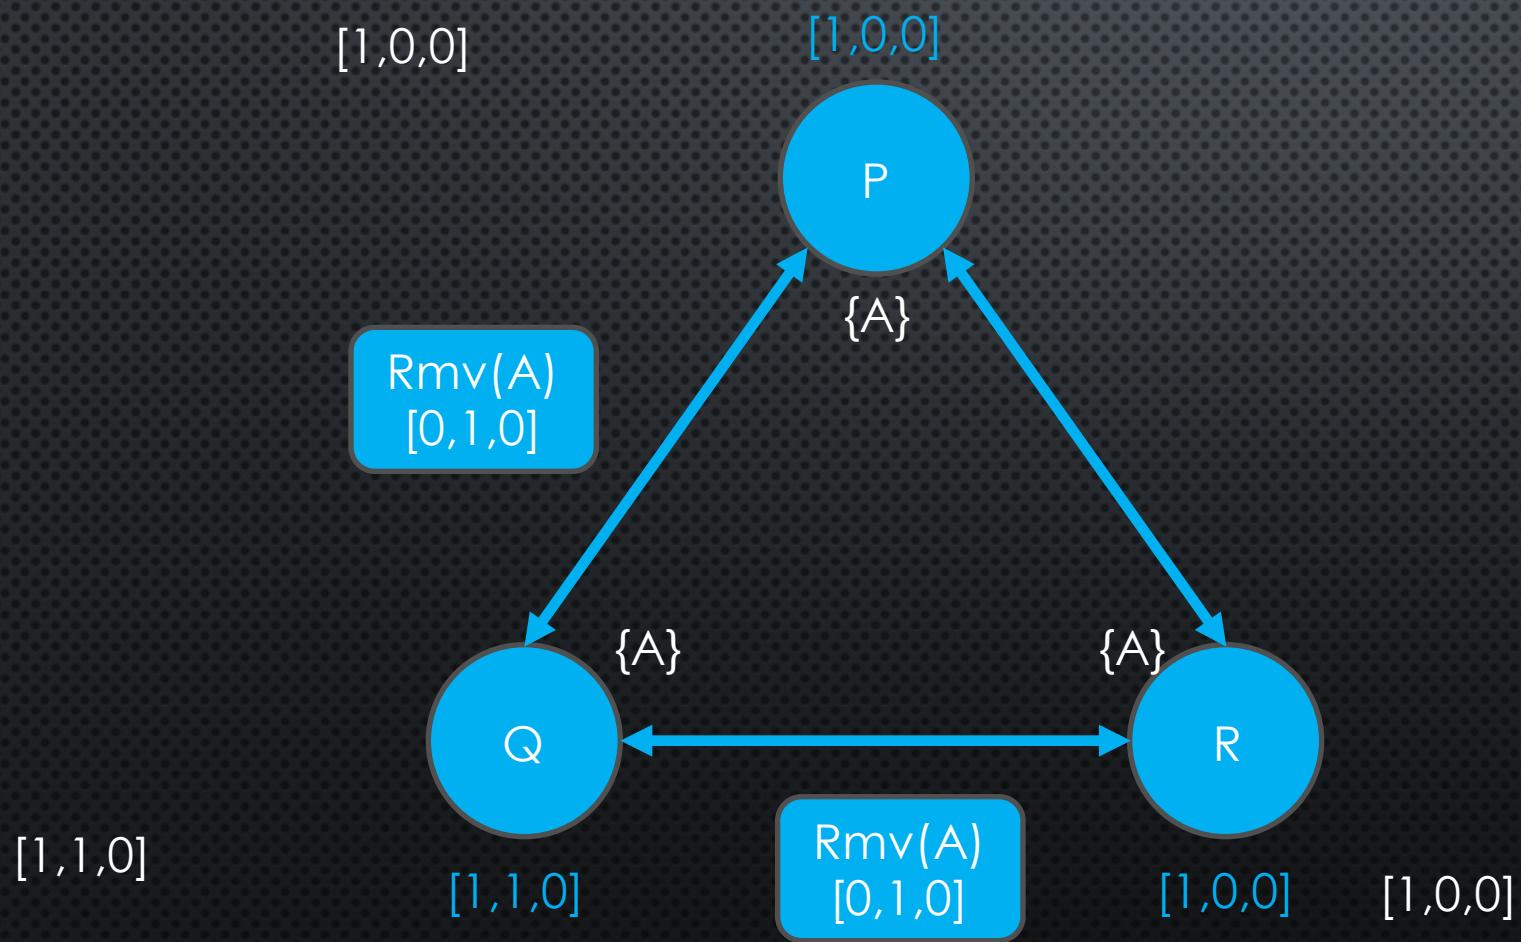
## SOLUTION 3



Scenario 2:

- Both ops tagged at App level
- P bcasts Add(A)
- Rmv(A) is still in the queue waiting to be bcast by Q
- All nodes deliver Add(A)

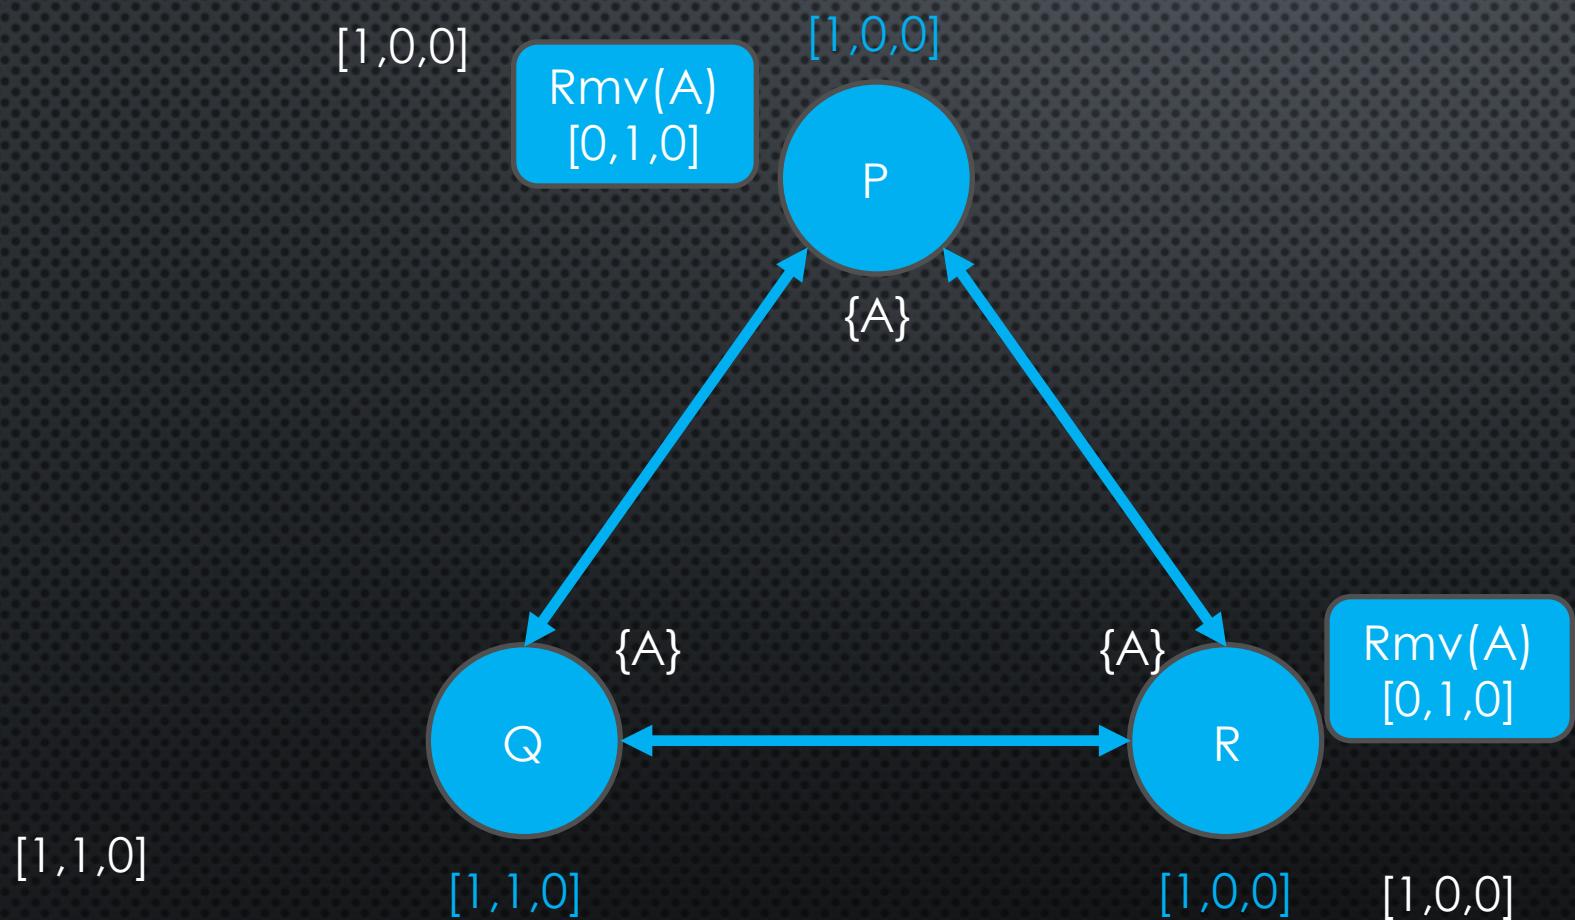
# SOLUTION 3



Scenario 2:

- Both ops tagged at App level
- P bcasts Add(A)
- Rmv(A) is still in the queue waiting to be bcast by Q
- All nodes deliver Add(A)
- Q dequeues Rmv(A) and bcasts it

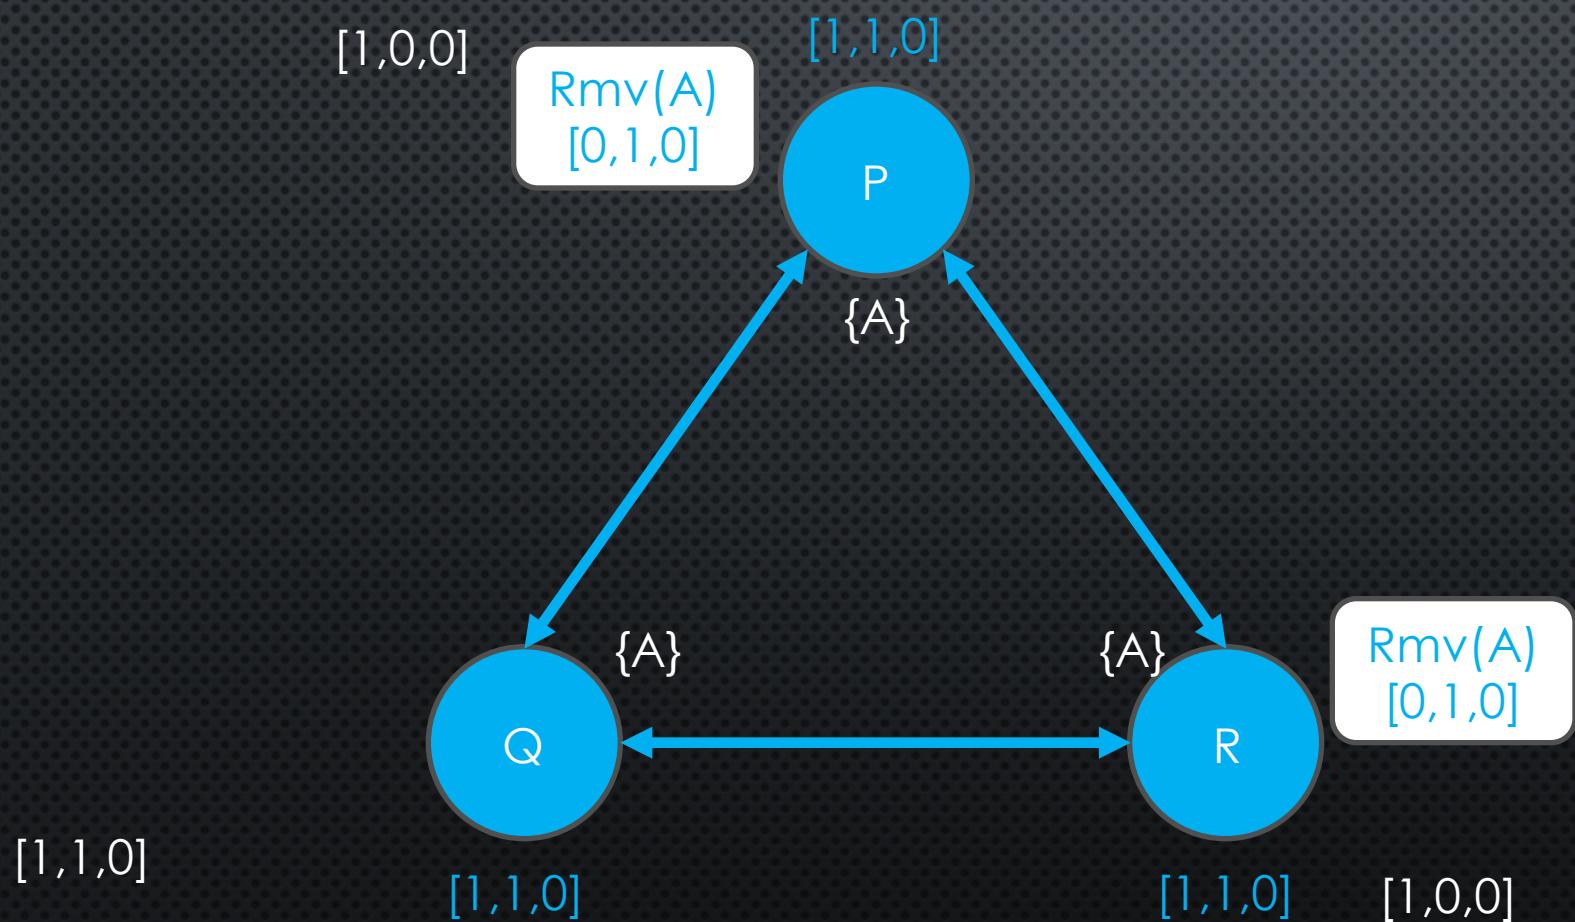
# SOLUTION 3



Scenario 2:

- Both ops tagged at App level
- P bcasts Add(A)
- Rmv(A) is still in the queue waiting to be bcast by Q
- All nodes deliver Add(A)
- Q dequeues Rmv(A) and bcasts it

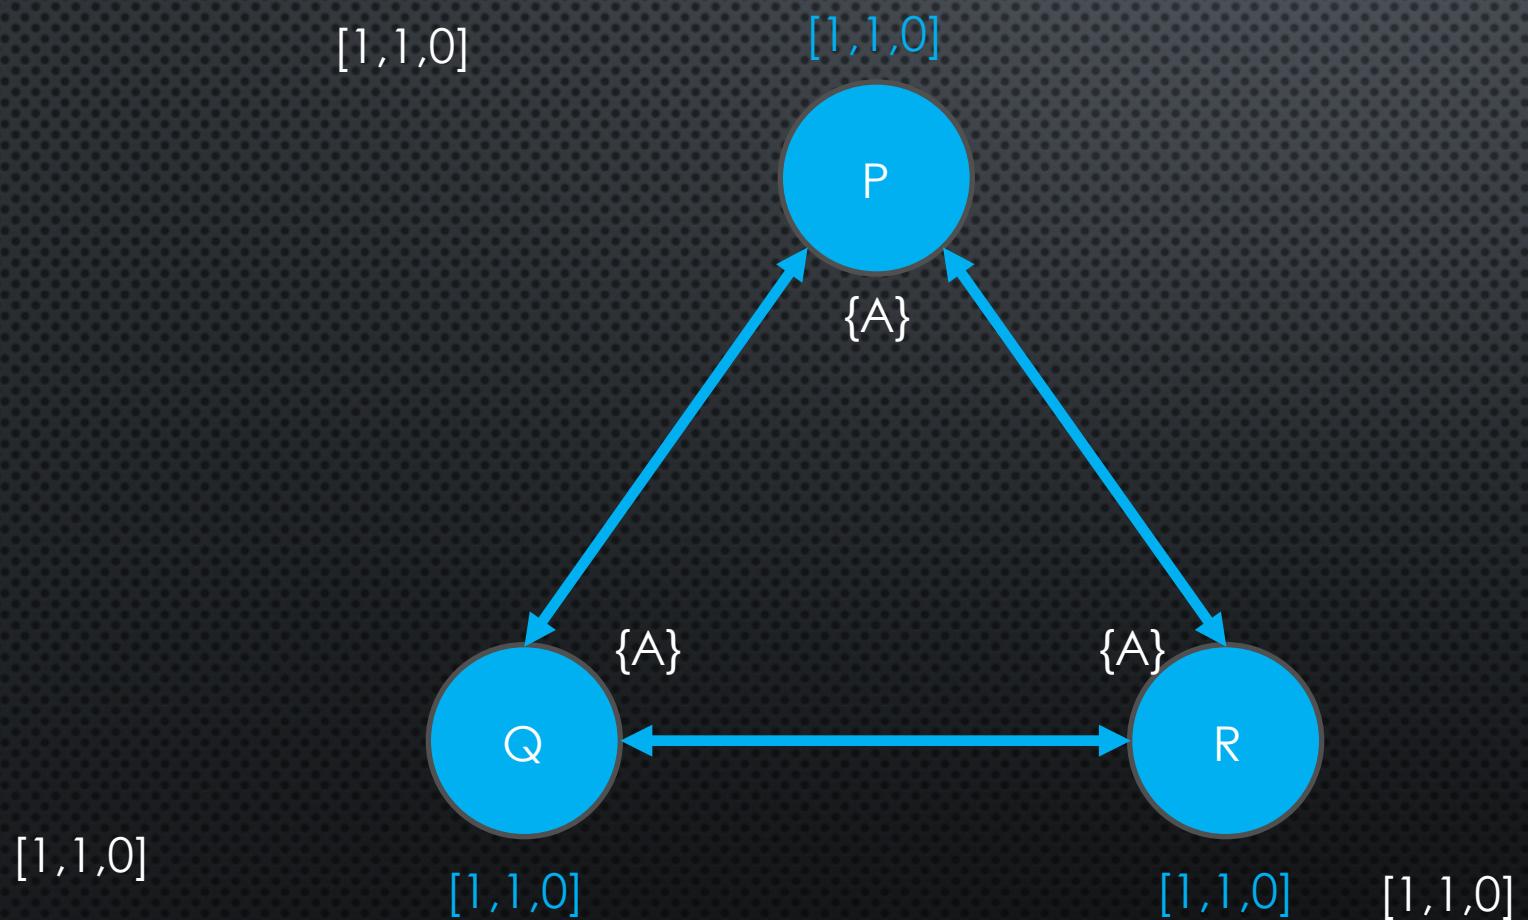
# SOLUTION 3



Scenario 2:

- Both ops tagged at App level
- P bcasts Add(A)
- Rmv(A) is still in the queue waiting to be bcast by Q
- All nodes deliver Add(A)
- Q dequeues Rmv(A) and bcasts it
- Rmv(A) delivered at P and R

## SOLUTION 3



### Scenario 2:

- Both ops tagged at App level
- P bcasts Add(A)
- Rmv(A) is still in the queue waiting to be bcast by Q
- All nodes deliver Add(A)
- Q dequeues Rmv(A) and bcasts it
- Rmv(A) delivered at P and R
- $[1,0,0]$  concurrent  $[0,1,0]$ 
  - Add Wins: A is there

# A WELCOME SIDE EFFECT

- AS YOU NOTICED IN SOLUTION 2 (SCENARIO 2):
  - CAUSAL DELIVERY SOMETIMES TAGS CONCURRENT OPERATIONS AS ONE HAPPENED-BEFORE THE OTHER
  - THIS LEADS TO OVER ORDERING OPERATIONS AND COULD INCLUDE EXTRA DELAY ON DELIVERY
  - THIS ALSO HAPPENS IN SOLUTION 1
  - THIS DOES NOT HAPPEN IN SOLUTION 3

# COMPARISON

	Solution 1	Solution 2	Solution 3
Tagging at Middleware	yes	yes	no
Tagging at Application	yes	no	yes
Characterization of Happened-before	correct	incorrect	correct
Over Ordering operations	yes	yes	no
Unneeded Slower Delivery	yes	yes	no
Solves the problem	yes	no	yes

# Questions ?

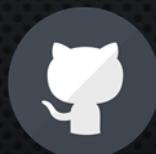
Slides: <https://bit.ly/tagged-causal>

Erlang implementation:

- Reliable Causal Broadcast: <https://github.com/gyounes/RCB>
- Tagged Reliable Causal Broadcast: [https://github.com/gyounes/trcb\\_base](https://github.com/gyounes/trcb_base)



@g\_unis



gyounes